

# CS336 Assignment 1 (basics): Building a Transformer LM

Version 1.0.4

CS336 Staff

Spring 2025

## 1 Assignment Overview

In this assignment, you will build all the components needed to train a standard Transformer language model (LM) from scratch and train some models.

### What you will implement

1. Byte-pair encoding (BPE) tokenizer (§2)
2. Transformer language model (LM) (§3)
3. The cross-entropy loss function and the AdamW optimizer (§4)
4. The training loop, with support for serializing and loading model and optimizer state (§5)

### What you will run

1. Train a BPE tokenizer on the TinyStories dataset.
2. Run your trained tokenizer on the dataset to convert it into a sequence of integer IDs.
3. Train a Transformer LM on the TinyStories dataset.
4. Generate samples and evaluate perplexity using the trained Transformer LM.
5. Train models on OpenWebText and submit your attained perplexities to a leaderboard.

**What you can use** We expect you to build these components from scratch. In particular, you may *not* use any definitions from `torch.nn`, `torch.nn.functional`, or `torch.optim` except for the following:

- `torch.nn.Parameter`
- Container classes in `torch.nn` (e.g., `Module`, `ModuleList`, `Sequential`, etc.)<sup>1</sup>
- The `torch.optim.Optimizer` base class

You may use any other PyTorch definitions. If you would like to use a function or class and are not sure whether it is permitted, feel free to ask on Slack. When in doubt, consider if using it compromises the “from-scratch” ethos of the assignment.

---

<sup>1</sup>See [PyTorch.org/docs/stable/nn.html#containers](https://pytorch.org/docs/stable/nn.html#containers) for a full list.

**Statement on AI tools** Prompting LLMs such as ChatGPT is permitted for low-level programming questions or high-level conceptual questions about language models, but using it directly to solve the problem is prohibited.

We strongly encourage you to disable AI autocomplete (e.g., Cursor Tab, GitHub CoPilot) in your IDE when completing assignments (though non-AI autocomplete, e.g., autocompleting function names is totally fine). We have found that AI autocomplete makes it much harder to engage deeply with the content.

**What the code looks like** All the assignment code as well as this writeup are available on GitHub at:

[github.com/stanford-cs336/assignment1-basics](https://github.com/stanford-cs336/assignment1-basics)

Please `git clone` the repository. If there are any updates, we will notify you so you can `git pull` to get the latest.

1. `cs336_basics/*`: This is where you write your code. Note that there’s no code in here—you can do whatever you want from scratch!
2. `adapters.py`: There is a set of functionality that your code must have. For each piece of functionality (e.g., scaled dot product attention), fill out its implementation (e.g., `run_scaled_dot_product_attention`) by simply invoking your code. Note: your changes to `adapters.py` should not contain any substantive logic; this is glue code.
3. `test_*.py`: This contains all the tests that you must pass (e.g., `test_scaled_dot_product_attention`), which will invoke the hooks defined in `adapters.py`. Don’t edit the test files.

**How to submit** You will submit the following files to Gradescope:

- `writeup.pdf`: Answer all the written questions. Please typeset your responses.
- `code.zip`: Contains all the code you’ve written.

To submit to the leaderboard, submit a PR to:

[github.com/stanford-cs336/assignment1-basics-leaderboard](https://github.com/stanford-cs336/assignment1-basics-leaderboard)

See the `README.md` in the leaderboard repository for detailed submission instructions.

**Where to get datasets** This assignment will use two pre-processed datasets: TinyStories [Eldan and Li, 2023] and OpenWebText [Gokaslan et al., 2019]. Both datasets are single, large plaintext files. If you are doing the assignment with the class, you can find these files at `/data` of any non-head node machine.

If you are following along at home, you can download these files with the commands inside the `README.md`.

### Low-Resource/Downscaling Tip: Init

Throughout the course’s assignment handouts, we will give advice for working through parts of the assignment with fewer or no GPU resources. For example, we will sometimes suggest **downscaling** your dataset or model size, or explain how to run training code on a MacOS integrated GPU or CPU. You’ll find these “low-resource tips” in a blue box (like this one). Even if you are an enrolled Stanford student with access to the course machines, these tips may help you iterate faster and save time, so we recommend you to read them!

### **Low-Resource/Downscaling Tip: Assignment 1 on Apple Silicon or CPU**

---

With the staff solution code, we can train an LM to generate reasonably fluent text on an Apple M3 Max chip with 36 GB RAM, in under 5 minutes on Metal GPU (MPS) and about 30 minutes using the CPU. If these words don't mean much to you, don't worry! Just know that if you have a reasonably up-to-date laptop and your implementation is correct and efficient, you will be able to train a small LM that generates simple children's stories with decent fluency.

Later in the assignment, we will explain what changes to make if you are on CPU or MPS.

## 2 Byte-Pair Encoding (BPE) Tokenizer

In the first part of the assignment, we will train and implement a byte-level byte-pair encoding (BPE) tokenizer [Sennrich et al., 2016, Wang et al., 2019]. In particular, we will represent arbitrary (Unicode) strings as a sequence of bytes and train our BPE tokenizer on this byte sequence. Later, we will use this tokenizer to encode text (a string) into tokens (a sequence of integers) for language modeling.

### 2.1 The Unicode Standard

Unicode is a text encoding standard that maps characters to integer *code points*. As of Unicode 16.0 (released in September 2024), the standard defines 154,998 characters across 168 scripts. For example, the character “s” has the code point 115 (typically notated as U+0073, where U+ is a conventional prefix and 0073 is 115 in hexadecimal), and the character “牛” has the code point 29275. In Python, you can use the `ord()` function to convert a single Unicode character into its integer representation. The `chr()` function converts an integer Unicode code point into a string with the corresponding character.

```
>>> ord('牛')
29275
>>> chr(29275)
'牛'
```

#### Problem (unicode1): Understanding Unicode (1 point)

- (a) What Unicode character does `chr(0)` return?

**Deliverable:** A one-sentence response.

- (b) How does this character’s string representation (`__repr__()`) differ from its printed representation?

**Deliverable:** A one-sentence response.

- (c) What happens when this character occurs in text? It may be helpful to play around with the following in your Python interpreter and see if it matches your expectations:

```
>>> chr(0)
>>> print(chr(0))
>>> "this is a test" + chr(0) + "string"
>>> print("this is a test" + chr(0) + "string")
```

**Deliverable:** A one-sentence response.

### 2.2 Unicode Encodings

While the Unicode standard defines a mapping from characters to code points (integers), it’s impractical to train tokenizers directly on Unicode codepoints, since the vocabulary would be prohibitively large (around 150K items) and sparse (since many characters are quite rare). Instead, we’ll use a Unicode encoding, which converts a Unicode character into a sequence of bytes. The Unicode standard itself defines three encodings: UTF-8, UTF-16, and UTF-32, with UTF-8 being the dominant encoding for the Internet (more than 98% of all webpages).

To encode a Unicode string into UTF-8, we can use the `encode()` function in Python. To access the underlying byte values for a Python `bytes` object, we can iterate over it (e.g., call `list()`). Finally, we can use the `decode()` function to decode a UTF-8 byte string into a Unicode string.

```

>>> test_string = "hello! こんにちは!"
>>> utf8_encoded = test_string.encode("utf-8")
>>> print(utf8_encoded)
b'hello! \xe3\x81\xe3\x82\xe3\x81\xab\xe3\x81\xa1\xe3\x81\xaf!'
>>> print(type(utf8_encoded))
<class 'bytes'>
>>> # Get the byte values for the encoded string (integers from 0 to 255).
>>> list(utf8_encoded)
[104, 101, 108, 108, 111, 33, 32, 227, 129, 147, 227, 130, 147, 227, 129, 171, 227, 129, 161, 227, 129, 175, 33]
>>> # One byte does not necessarily correspond to one Unicode character!
>>> print(len(test_string))
13
>>> print(len(utf8_encoded))
23
>>> print(utf8_encoded.decode("utf-8"))
hello! こんにちは!

```

By converting our Unicode codepoints into a sequence of bytes (e.g., via the UTF-8 encoding), we are essentially taking a sequence of codepoints (integers in the range 0 to 154,997) and transforming it into a sequence of byte values (integers in the range 0 to 255). The 256-length byte vocabulary is *much* more manageable to deal with. When using byte-level tokenization, we do not need to worry about out-of-vocabulary tokens, since we know that *any* input text can be expressed as a sequence of integers from 0 to 255.

#### Problem (unicode2): Unicode Encodings (3 points)

- (a) What are some reasons to prefer training our tokenizer on UTF-8 encoded bytes, rather than UTF-16 or UTF-32? It may be helpful to compare the output of these encodings for various input strings.

**Deliverable:** A one-to-two sentence response.

- (b) Consider the following (incorrect) function, which is intended to decode a UTF-8 byte string into a Unicode string. Why is this function incorrect? Provide an example of an input byte string that yields incorrect results.

bytes with an integer return a bytestring of that integer length,  
while bytes with an iterable return the bytestring of that object

```

def decode_utf8_bytes_to_str_wrong(bytestring: bytes):
    return "".join([bytes([b]).decode("utf-8") for b in bytestring])

```

```

>>> decode_utf8_bytes_to_str_wrong("hello".encode("utf-8"))
'hello'

```

**Deliverable:** An example input byte string for which `decode_utf8_bytes_to_str_wrong` produces incorrect output, with a one-sentence explanation of why the function is incorrect.

- (c) Give a two byte sequence that does not decode to any Unicode character(s).

**Deliverable:** An example, with a one-sentence explanation.

## 2.3 Subword Tokenization

While byte-level tokenization can alleviate the out-of-vocabulary issues faced by word-level tokenizers, tokenizing text into bytes results in extremely long input sequences. This slows down model training, since a

sentence with 10 words might only be 10 tokens long in a word-level language model, but could be 50 or more tokens long in a character-level model (depending on the length of the words). Processing these longer sequences requires more computation at each step of the model. Furthermore, **language modeling on byte sequences is difficult because the longer input sequences create long-term dependencies in the data.**

Subword tokenization is a midpoint between word-level tokenizers and byte-level tokenizers. Note that a byte-level tokenizer’s vocabulary has 256 entries (byte values are 0 to 225). A subword tokenizer trades-off a larger vocabulary size for better compression of the input byte sequence. For example, if the byte sequence `b'the'` often occurs in our raw text training data, assigning it an entry in the vocabulary would reduce this 3-token sequence to a single token.

How do we select these subword units to add to our vocabulary? Sennrich et al. [2016] propose to use byte-pair encoding (BPE; Gage, 1994), a compression algorithm that iteratively replaces (“merges”) the most frequent pair of bytes with a single, new unused index. Note that this algorithm adds subword tokens to our vocabulary to maximize the compression of our input sequences—if a word occurs in our input text enough times, it’ll be represented as a single subword unit.

Subword tokenizers with vocabularies constructed via BPE are often called BPE tokenizers. In this assignment, we’ll implement a byte-level BPE tokenizer, where the vocabulary items are bytes or merged sequences of bytes, which give us the best of both worlds in terms of out-of-vocabulary handling and manageable input sequence lengths. **The process of constructing the BPE tokenizer vocabulary is known as “training” the BPE tokenizer.**

## 2.4 BPE Tokenizer Training

The BPE tokenizer training procedure consists of three main steps.

**Vocabulary initialization** The tokenizer vocabulary is a one-to-one mapping from `bytestring` token to integer ID. Since we’re training a byte-level BPE tokenizer, our initial vocabulary is simply the set of all bytes. Since there are 256 possible byte values, our initial vocabulary is of size 256.

**Pre-tokenization** Once you have a vocabulary, you could, in principle, count how often bytes occur next to each other in your text and begin merging them starting with the most frequent pair of bytes. However, this is quite computationally expensive, since we’d have to go take a full pass over the corpus each time we merge. **In addition, directly merging bytes across the corpus may result in tokens that differ only in punctuation (e.g., `dog!` vs. `dog.`).** These tokens would get completely different token IDs, even though they are likely to have high semantic similarity (since they differ only in punctuation).

To avoid this, we **pre-tokenize** the corpus. You can think of this as a coarse-grained tokenization over the corpus that helps us count how often pairs of characters appear. For example, the word `'text'` might be a pre-token that appears 10 times. In this case, when we count how often the characters `'t'` and `'e'` appear next to each other, we will see that the word `'text'` has `'t'` and `'e'` adjacent and we can increment their count by 10 instead of looking through the corpus. Since we’re training a byte-level BPE model, each pre-token is represented as a sequence of UTF-8 bytes.

The original BPE implementation of Sennrich et al. [2016] pre-tokenizes by simply splitting on whitespace (i.e., `s.split(" ")`). In contrast, we’ll use a regex-based pre-tokenizer (used by GPT-2; Radford et al., 2019) from [github.com/openai/tiktoken/pull/234/files](https://github.com/openai/tiktoken/pull/234/files):

```
>>> PAT = r'(?[sdmt]|ll|ve|re) ?\p{L}+ ?\p{N}+ ?[^\s\p{L}\p{N}]+\s+(?!S)|\s+'''
```

It may be useful to interactively split some text with this pre-tokenizer to get a better sense of its behavior:

```
>>> # requires `regex` package
>>> import regex as re
>>> re.findall(PAT, "some text that i'll pre-tokenize")
['some', ' ' text', ' ' that', ' ' i', ' 'll', ' ' pre', ' -', 'tokenize']
```

When using it in your code, however, you should use `re.finditer` to avoid storing the pre-tokenized words as you construct your mapping from pre-tokens to their counts.

**Compute BPE merges** Now that we’ve converted our input text into pre-tokens and represented each pre-token as a sequence of UTF-8 bytes, we can compute the BPE merges (i.e., train the BPE tokenizer). At a high level, the BPE algorithm iteratively counts every pair of bytes and identifies the pair with the highest frequency (“A”, “B”). Every occurrence of this most frequent pair (“A”, “B”) is then *merged*, i.e., replaced with a new token “AB”. This new merged token is added to our vocabulary; as a result, the final vocabulary after BPE training is the size of the initial vocabulary (256 in our case), plus the number of BPE merge operations performed during training. For efficiency during BPE training, we do not consider pairs that cross pre-token boundaries.<sup>2</sup> When computing merges, deterministically break ties in pair frequency by *preferring the lexicographically greater pair*. For example, if the pairs (“A”, “B”), (“A”, “C”), (“B”, “ZZ”), and (“BA”, “A”) all have the highest frequency, we’d merge (“BA”, “A”):

```
>>> max([("A", "B"), ("A", "C"), ("B", "ZZ"), ("BA", "A")])
('BA', 'A')
```

**Special tokens** Often, some strings (e.g., `<|endoftext|>`) are used to encode metadata (e.g., boundaries between documents). When encoding text, it’s often desirable to treat some strings as “special tokens” that should never be split into multiple tokens (i.e., will always be preserved as a single token). For example, the end-of-sequence string `<|endoftext|>` should always be preserved as a single token (i.e., a single integer ID), so we know when to stop generating from the language model. These special tokens must be added to the vocabulary, so they have a corresponding fixed token ID.

Algorithm 1 of Sennrich et al. [2016] contains an inefficient implementation of BPE tokenizer training (essentially following the steps that we outlined above). As a first exercise, it may be useful to implement and test this function to test your understanding.

#### Example (bpe\_example): BPE training example

Here is a stylized example from Sennrich et al. [2016]. Consider a corpus consisting of the following text

```
low low low low low
lower lower widest widest widest
newest newest newest newest newest newest
```

and the vocabulary has a special token `<|endoftext|>`.

**Vocabulary** We initialize our vocabulary with our special token `<|endoftext|>` and the 256 byte values.

**Pre-tokenization** For simplicity and to focus on the merge procedure, we assume in this example that pretokenization simply splits on whitespace. When we pretokenize and count, we end up with the frequency table.

```
{low: 5, lower: 2, widest: 3, newest: 6}
```

<sup>2</sup>Note that the original BPE formulation [Sennrich et al., 2016] specifies the inclusion of an end-of-word token. We do not add an end-of-word-token when training byte-level BPE models because all bytes (including whitespace and punctuation) are included in the model’s vocabulary. Since we’re explicitly representing spaces and punctuation, the learned BPE merges will naturally reflect these word boundaries.

It is convenient to represent this as a `dict[tuple[bytes], int]`, e.g. `{(1,0,w): 5 ...}`. Note that even a single byte is a `bytes` object in Python. **There is no `byte` type in Python to represent a single byte, just as there is no `char` type in Python to represent a single character.**

**Merges** We first look at every successive pair of bytes and sum the frequency of the words where they appear `{lo: 7, ow: 7, we: 8, er: 2, wi: 3, id: 3, de: 3, es: 9, st: 9, ne: 6, ew: 6}`. The pair `('es')` and `('st')` are tied, so we take the lexicographically greater pair, `('st')`. We would then merge the pre-tokens so that we end up with `{(1,0,w): 5, (1,0,w,e,r): 2, (w,i,d,e,st): 3, (n,e,w,e,st): 6}`.

In the second round, we see that `(e, st)` is the most common pair (with a count of 9) and we would merge into `{(1,0,w): 5, (1,0,w,e,r): 2, (w,i,d,est): 3, (n,e,w,est): 6}`. Continuing this, the sequence of merges we get in the end will be `['s t', 'e st', 'o w', 'l ow', 'w est', 'n e', 'ne west', 'w i', 'wi d', 'wid est', 'low e', 'lowe r']`.

If we take 6 merges, we have `['s t', 'e st', 'o w', 'l ow', 'w est', 'n e']` and our vocabulary elements would be `<|endoftext|>, [...256 BYTE CHARS], st, est, ow, low, west, ne`.

With this vocabulary and set of merges, the word `newest` would tokenize as `[ne, west]`.

## 2.5 Experimenting with BPE Tokenizer Training

Let's train a byte-level BPE tokenizer on the TinyStories dataset. Instructions to find / download the dataset can be found in Section 1. Before you start, we recommend taking a look at the TinyStories dataset to get a sense of what's in the data.

**Parallelizing pre-tokenization** You will find that a major bottleneck is the pre-tokenization step. You can speed up pre-tokenization by parallelizing your code with the built-in library `multiprocessing`. Concretely, we recommend that in parallel implementations of pre-tokenization, you chunk the corpus while ensuring your chunk boundaries occur at the beginning of a special token. You are free to use the starter code at the following link verbatim to obtain chunk boundaries, which you can then use to distribute work across your processes:

[https://github.com/stanford-cs336/assignment1-basics/blob/main/cs336\\_basics/pretokenization\\_example.py](https://github.com/stanford-cs336/assignment1-basics/blob/main/cs336_basics/pretokenization_example.py)

This chunking will always be valid, since we never want to merge across document boundaries. For the purposes of the assignment, you can always split in this way. **Don't worry about the edge case of receiving a very large corpus that does not contain `<|endoftext|>`.**

**Removing special tokens before pre-tokenization** Before running pre-tokenization with the regex pattern (using `re.finditer`), you should strip out all special tokens from your corpus (or your chunk, if using a parallel implementation). Make sure that you **split** on your special tokens, so that no merging can occur across the text they delimit. For example, if you have a corpus (or chunk) like `[Doc 1]<|endoftext|>[Doc 2]`, you should split on the special token `<|endoftext|>`, and pre-tokenize `[Doc 1]` and `[Doc 2]` separately, so that no merging can occur across the document boundary. This can be done using `re.split` with `"|"`. `.join(special_tokens)` as the delimiter (with careful use of `re.escape` since `|` may occur in the special tokens). The test `test_train_bpe_special_tokens` will test for this.

**Optimizing the merging step** The naïve implementation of BPE training in the stylized example above is slow because for every merge, it iterates over all byte pairs to identify the most frequent pair. However, the only pair counts that change after each merge are those that overlap with the merged pair. Thus, BPE training speed can be improved by indexing the counts of all pairs and incrementally updating these counts, rather than explicitly iterating over each pair of bytes to count pair frequencies. You can get significant speedups with this caching procedure, though we note that the merging part of BPE training is *not* parallelizable in Python.



### Low-Resource/Downscaling Tip: Profiling

You should use profiling tools like `cProfile` or `scalene` to identify the bottlenecks in your implementation, and focus on optimizing those.

### Low-Resource/Downscaling Tip: “Downscaling”

Instead of jumping to training your tokenizer on the full TinyStories dataset, we recommend you first train on a small subset of the data: a “debug dataset”. For example, you could train your tokenizer on the TinyStories validation set instead, which is 22K documents instead of 2.12M. This illustrates a general strategy of downscaling whenever possible to speed up development: for example, using smaller datasets, smaller model sizes, etc. Choosing the size of the debug dataset or hyperparameter config requires careful consideration: you want your debug set to be large enough to have the same bottlenecks as the full configuration (so that the optimizations you make will generalize), but not so big that it takes forever to run.

### Problem (train\_bpe): BPE Tokenizer Training (15 points)

**Deliverable:** Write a function that, given a path to an input text file, trains a (byte-level) BPE tokenizer. Your BPE training function should handle (at least) the following input parameters:

**input\_path:** `str` Path to a text file with BPE tokenizer training data.

**vocab\_size:** `int` A positive integer that defines the maximum final vocabulary size (including the initial byte vocabulary, vocabulary items produced from merging, and any special tokens).

**special\_tokens:** `list[str]` A list of strings to add to the vocabulary. These special tokens do not otherwise affect BPE training.

Your BPE training function should return the resulting vocabulary and merges:

**vocab:** `dict[int, bytes]` The tokenizer vocabulary, a mapping from `int` (token ID in the vocabulary) to `bytes` (token bytes).

**merges:** `list[tuple[bytes, bytes]]` A list of BPE merges produced from training. Each list item is a `tuple` of `bytes` (`<token1>`, `<token2>`), representing that `<token1>` was merged with `<token2>`. The merges should be ordered by order of creation.

To test your BPE training function against our provided tests, you will first need to implement the test adapter at `[adapters.run_train_bpe]`. Then, run `uv run pytest tests/test_train_bpe.py`. Your implementation should be able to pass all tests. Optionally (this could be a large time-investment), you can implement the key parts of your training method using some systems language, for instance C++ (consider `cppyy` for this) or Rust (using `PyO3`). If you do this, be aware of which operations require copying vs reading directly from Python memory, and make sure to leave build instructions, or make sure it builds using only `pyproject.toml`. Also note that the GPT-2 regex is not well-supported in most regex engines and will be too slow in most that do. We have verified that `Oniguruma` is reasonably fast and supports negative lookahead, but the `regex` package in Python is, if anything, even faster.

### Problem (train\_bpe\_tinystories): BPE Training on TinyStories (2 points)

- (a) Train a byte-level BPE tokenizer on the TinyStories dataset, using a maximum vocabulary size of 10,000. Make sure to add the TinyStories `<|endoftext|>` special token to the vocabulary. Serialize the resulting vocabulary and merges to disk for further inspection. How many hours and memory did training take? What is the longest token in the vocabulary? Does it make sense?

**Resource requirements:**  $\leq 30$  minutes (no GPUs),  $\leq 30$ GB RAM

**Hint** You should be able to get under 2 minutes for BPE training using `multiprocessing` during pretokenization and the following two facts:

- (a) The `<|endoftext|>` token delimits documents in the data files.
- (b) The `<|endoftext|>` token is handled as a special case before the BPE merges are applied.

**Deliverable:** A one-to-two sentence response.

- (b) Profile your code. What part of the tokenizer training process takes the most time?

**Deliverable:** A one-to-two sentence response.

Next, we'll try training a byte-level BPE tokenizer on the OpenWebText dataset. As before, we recommend taking a look at the dataset to better understand its contents.

### Problem (train\_bpe\_expts\_owt): BPE Training on OpenWebText (2 points)

- (a) Train a byte-level BPE tokenizer on the OpenWebText dataset, using a maximum vocabulary size of 32,000. Serialize the resulting vocabulary and merges to disk for further inspection. What is the longest token in the vocabulary? Does it make sense?

**Resource requirements:**  $\leq 12$  hours (no GPUs),  $\leq 100$ GB RAM

**Deliverable:** A one-to-two sentence response.

- (b) Compare and contrast the tokenizer that you get training on TinyStories versus OpenWebText.

**Deliverable:** A one-to-two sentence response.

## 2.6 BPE Tokenizer: Encoding and Decoding

In the previous part of the assignment, we implemented a function to train a BPE tokenizer on input text to obtain a tokenizer vocabulary and a list of BPE merges. Now, we will implement a BPE tokenizer that loads a provided vocabulary and list of merges and uses them to encode and decode text to/from token IDs.

### 2.6.1 Encoding text

The process of encoding text by BPE mirrors how we train the BPE vocabulary. There are a few major steps.

**Step 1: Pre-tokenize.** We first pre-tokenize the sequence and represent each pre-token as a sequence of UTF-8 bytes, just as we did in BPE training. We will be merging these bytes within each pre-token into vocabulary elements, handling each pre-token independently (no merges across pre-token boundaries).

**Step 2: Apply the merges.** We then take the sequence of vocabulary element merges created during BPE training, and apply it to our pre-tokens *in the same order of creation*.

### Example (bpe\_encoding): BPE encoding example

For example, suppose our input string is 'the cat ate', our vocabulary is {0: b' ', 1: b'a', 2: b'c', 3: b'e', 4: b'h', 5: b't', 6: b'th', 7: b' c', 8: b' a', 9: b'the', 10: b' at'}, and our learned merges are [(b't', b'h'), (b' ', b'c'), (b' ', 'a'), (b'th', b'e'), (b' a', b't')]. First, our pre-tokenizer would split this string into ['the', ' cat', ' ate']. Then, we'll look at each pre-token and apply the BPE merges.

The first pre-token 'the' is initially represented as [b't', b'h', b'e']. Looking at our list of merges, we identify the first applicable merge to be (b't', b'h'), and use that to transform the pre-token into [b'th', b'e']. Then, we go back to the list of merges and identify the next applicable merge to be (b'th', b'e'), which transforms the pre-token into [b'the']. Finally, looking back at the list of merges, we see that there are no more that apply to the string (since the entire pre-token has been merged into a single token), so we are done applying the BPE merges. The corresponding integer sequence is [9].

Repeating this process for the remaining pre-tokens, we see that the pre-token ' cat' is represented as [b' c', b'a', b't'] after applying the BPE merges, which becomes the integer sequence [7, 1, 5]. The final pre-token ' ate' is [b' at', b'e'] after applying the BPE merges, which becomes the integer sequence [10, 3]. Thus, the final result of encoding our input string is [9, 7, 1, 5, 10, 3].

**Special tokens.** Your tokenizer should be able to properly handle user-defined special tokens when encoding text (provided when constructing the tokenizer).

**Memory considerations.** Suppose we want to tokenize a large text file that we cannot fit in memory. To efficiently tokenize this large file (or any other stream of data), we need to break it up into manageable chunks and process each chunk in-turn, so that the memory complexity is constant as opposed to linear in the size of the text. In doing so, we need to make sure that a token doesn't cross chunk boundaries, else we'll get a different tokenization than the naïve method of tokenizing the entire sequence in-memory.

### 2.6.2 Decoding text

To decode a sequence of integer token IDs back to raw text, we can simply look up each ID's corresponding entries in the vocabulary (a byte sequence), concatenate them together, and then decode the bytes to a Unicode string. Note that input IDs are not guaranteed to map to valid Unicode strings (since a user could input any sequence of integer IDs). In the case that the input token IDs do not produce a valid Unicode string, you should replace the malformed bytes with the official Unicode replacement character U+FFFD.<sup>3</sup> The `errors` argument of `bytes.decode` controls how Unicode decoding errors are handled, and using `errors='replace'` will automatically replace malformed data with the replacement marker.

### Problem (tokenizer): Implementing the tokenizer (15 points)

**Deliverable:** Implement a `Tokenizer` class that, given a vocabulary and a list of merges, encodes text into integer IDs and decodes integer IDs into text. Your tokenizer should also support user-provided special tokens (appending them to the vocabulary if they aren't already there). We recommend the following interface:

```
def __init__(self, vocab, merges, special_tokens=None) Construct a tokenizer from a given
    vocabulary, list of merges, and (optionally) a list of special tokens. This function should accept
```

<sup>3</sup>See [en.wikipedia.org/wiki/Specials\\_\(Unicode\\_block\)#Replacement\\_character](https://en.wikipedia.org/wiki/Specials_(Unicode_block)#Replacement_character) for more information about the Unicode replacement character.

the following parameters:

```
vocab: dict[int, bytes]
merges: list[tuple[bytes, bytes]]
special_tokens: list[str] | None = None
```

```
def from_files(cls, vocab_filepath, merges_filepath, special_tokens=None) Class
method that constructs and return a Tokenizer from a serialized vocabulary and list of merges
(in the same format that your BPE training code output) and (optionally) a list of special
tokens. This method should accept the following additional parameters:
```

```
vocab_filepath: str
merges_filepath: str
special_tokens: list[str] | None = None
```

```
def encode(self, text: str) -> list[int] Encode an input text into a sequence of token IDs.
```

```
def encode_iterable(self, iterable: Iterable[str]) -> Iterator[int] Given an iterable of
strings (e.g., a Python file handle), return a generator that lazily yields token IDs. This is
required for memory-efficient tokenization of large files that we cannot directly load into
memory.
```

```
def decode(self, ids: list[int]) -> str Decode a sequence of token IDs into text.
```

To test your `Tokenizer` against our provided tests, you will first need to implement the test adapter at `[adapters.get_tokenizer]`. Then, run `uv run pytest tests/test_tokenizer.py`. Your implementation should be able to pass all tests.

## 2.7 Experiments

### Problem (tokenizer\_experiments): Experiments with tokenizers (4 points)

- (a) Sample 10 documents from TinyStories and OpenWebText. Using your previously-trained TinyStories and OpenWebText tokenizers (10K and 32K vocabulary size, respectively), encode these sampled documents into integer IDs. What is each tokenizer's compression ratio (bytes/token)?

**Deliverable:** A one-to-two sentence response.

- (b) What happens if you tokenize your OpenWebText sample with the TinyStories tokenizer? Compare the compression ratio and/or qualitatively describe what happens.

**Deliverable:** A one-to-two sentence response.

- (c) Estimate the throughput of your tokenizer (e.g., in bytes/second). How long would it take to tokenize the Pile dataset (825GB of text)?

**Deliverable:** A one-to-two sentence response.

- (d) Using your TinyStories and OpenWebText tokenizers, encode the respective training and development datasets into a sequence of integer token IDs. We'll use this later to train our language model. We recommend serializing the token IDs as a NumPy array of datatype `uint16`. Why is `uint16` an appropriate choice?

**Deliverable:** A one-to-two sentence response.

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snyder/cs336\_basics/train\_bpe.py

```

1  import os
2  import heapq
3  from typing import BinaryIO
4  import regex as re
5  import collections
6  import multiprocessing as mp
7  import time
8  import pickle
9  from functools import reduce
10
11 # Regex for coarse tokenization
12 PAT = re.compile(r'([?:[sdmt]|ll|ve|re)| ?\p{L}+| ?\p{N}+| ?(?:\s\p{L}\p{N}+)|\s+
13      (?!S)|\s+)'
14
15 class ReverseLexOrderPair:
16     """
17     Encapsulates (bytes, bytes) so that in a min-heap, the "largest in normal lex
18     order"
19     is treated as the smallest. Ensures that tie frequencies pop in reverse lex
20     order.
21     """
22
23     def __init__(self, pair: tuple[bytes, bytes]):
24         self.pair = pair
25
26     def __lt__(self, other: "ReverseLexOrderPair") -> bool:
27         # Invert normal order: self < other if self is > other (so larger lex sorts
28         first).
29         return self.pair > other.pair
30
31     def __eq__(self, other: "ReverseLexOrderPair") -> bool:
32         return self.pair == other.pair
33
34 def find_chunk_boundaries(file: BinaryIO, desired_num_chunks: int,
35     split_special_token: bytes) -> list[int]:
36     """
37     Find chunk boundaries by reading forward from guessed positions
38     until split_special_token is found (or EOF). Ensures alignment.
39     """
40     assert isinstance(split_special_token, bytes), "Must represent special token as a
41     bytestring"
42
43     file.seek(0, os.SEEK_END)
44     file_size = file.tell()
45     file.seek(0)
46     chunk_size = file_size // desired_num_chunks

```

```

43
44     # Initial boundary guesses (uniformly spaced); force last boundary at EOF
45     chunk_boundaries = [i * chunk_size for i in range(desired_num_chunks + 1)]
46     chunk_boundaries[-1] = file_size
47
48     # This seems inefficient since it does extra work when the special token is not
49     # found in the mini chunk...
50     # but I guess we never expect this to fail
51     mini_chunk_size = 4096
52     for bi in range(1, len(chunk_boundaries) - 1):
53         pos = chunk_boundaries[bi]
54         file.seek(pos)
55         while True:
56             mini_chunk = file.read(mini_chunk_size)
57             if not mini_chunk:
58                 # If EOF is reached before finding split token
59                 chunk_boundaries[bi] = file_size
60                 break
61             found_at = mini_chunk.find(split_special_token)
62             if found_at != -1:
63                 # Found the split token; adjust boundary precisely
64                 chunk_boundaries[bi] = pos + found_at
65                 break
66             pos += mini_chunk_size
67
68     return sorted(set(chunk_boundaries)) #sorted turns a set -> list (that is sorted)
69
70 def pre_tokenize_chunk(chunk: str, special_pattern: re.Pattern | None) ->
71 dict[tuple[bytes], int]:
72     """Regex tokenizes the chunk. Splits first on special tokens, then uses PAT."""
73     freqs: dict[tuple[bytes], int] = {}
74     sub_chunks = special_pattern.split(chunk) if special_pattern else [chunk]
75
76     for sub_chunk in sub_chunks:
77         for match in PAT.finditer(sub_chunk):
78             match_bytes = tuple(bytes([b]) for b in match.group().encode("UTF-8"))
79             freqs[match_bytes] = freqs.get(match_bytes, 0) + 1
80
81     return freqs
82
83 def merge_freq_dicts(dict1: dict[tuple[bytes], int], dict2: dict[tuple[bytes], int])
84 -> dict[tuple[bytes], int]:
85     """Adds frequencies from dict2 into dict1."""
86     result = dict1.copy()
87     for key, value in dict2.items():
88         result[key] = result.get(key, 0) + value
89     return result
90

```

```

91 def pre_tokenize(input_path: str, special_tokens: list[str]) -> dict[tuple[bytes],
int]:
92     """
93     Splits a file into chunks aligned with <|endoftext|>, then tokenizes each chunk
94     in parallel. Returns aggregated frequency dict.
95     """
96     num_processes = mp.cpu_count()
97     pool = mp.Pool(processes=num_processes)
98     chunk_freqs = []
99     special_pattern = re.compile("|".join(re.escape(tok) for tok in special_tokens))
100 if special_tokens else None
101
102     with open(input_path, "rb") as f:
103         # Divide file into number of chunks matching number of cpu cores
104         boundaries = find_chunk_boundaries(f, num_processes, b"<|endoftext|>")
105
106         # Read each chunk in bytes, decode, then apply_async for parallel
107         # tokenization
108         for start, end in zip(boundaries[:-1], boundaries[1:]):
109             f.seek(start)
110             chunk_bytes = f.read(end - start)
111             chunk_str = chunk_bytes.decode("utf-8", errors="ignore")
112             chunk_freqs.append(pool.apply_async(pre_tokenize_chunk, (chunk_str,
113             special_pattern)))
114
115     pool.close()
116     pool.join()
117
118     # Collect and merge partial results
119     freq_dicts = [res.get() for res in chunk_freqs]
120     combined_freqs = reduce(merge_freq_dicts, freq_dicts, {})
121     return combined_freqs
122
123 def get_pair_freqs(
124     freqs: dict[tuple[bytes], int],
125 ) -> tuple[dict[tuple[bytes, bytes], int], dict[tuple[bytes, bytes],
126 set[tuple[bytes]]]]:
127     """
128     Builds a pair-frequency table and reverse mapping (pair -> set of keys).
129     """
130     pair_freqs: dict[tuple[bytes, bytes], int] = collections.defaultdict(int)
131     pairs_to_keys: dict[tuple[bytes, bytes], set[tuple[bytes]]] =
132     collections.defaultdict(set)
133
134     for symbols, freq in freqs.items():
135         for i in range(len(symbols) - 1):
136             pair = (symbols[i], symbols[i + 1])
137             pair_freqs[pair] += freq
138             pairs_to_keys[pair].add(symbols)

```



```

136     return pair_freqs, pairs_to_keys
137
138
139 def build_new_repr(old_repr: tuple[bytes], pair: tuple[bytes, bytes]) ->
tuple[bytes]:
140     """Replaces every occurrence of pair=(x,y) in old_repr with the merged symbol
x+y."""
141     new_symbols = []
142     i = 0
143     while i < len(old_repr):
144         if i < len(old_repr) - 1 and old_repr[i] == pair[0] and old_repr[i + 1] ==
pair[1]:
145             new_symbols.append(old_repr[i] + old_repr[i + 1]) # merges, e.g. b'A' +
b'B' => b'AB'
146             i += 2
147         else:
148             new_symbols.append(old_repr[i])
149             i += 1
150     return tuple(new_symbols)
151
152
153 def merge(
154     freqs: dict[tuple[bytes], int],
155     pair_freqs: dict[tuple[bytes, bytes], int],
156     pairs_to_keys: dict[tuple[bytes, bytes], set[tuple[bytes]]],
157     pair: tuple[bytes, bytes],
158 ) -> set[tuple[bytes, bytes]]:
159     """Merges 'pair' into freqs and updates pair_freqs & pairs_to_keys for all
affected old/new keys."""
160     changed_pairs = set()
161     keys_to_modify = pairs_to_keys[pair].copy()
162
163     for old_key in keys_to_modify:
164         old_freq = freqs.pop(old_key)
165         new_key = build_new_repr(old_key, pair)
166
167         # Decrement frequencies in pair_freqs for old_key's adjacencies
168         for i in range(len(old_key) - 1):
169             left, right = old_key[i], old_key[i + 1]
170             pair_freqs[left, right] -= old_freq
171             changed_pairs.add((left, right))
172             if pair_freqs[left, right] <= 0:
173                 del pair_freqs[left, right]
174             pairs_to_keys[left, right].discard(old_key)
175
176         # Increment frequencies for new_key's adjacencies
177         for i in range(len(new_key) - 1):
178             left, right = new_key[i], new_key[i + 1]
179             pair_freqs[left, right] += old_freq
180             changed_pairs.add((left, right))
181             pairs_to_keys[left, right].add(new_key)

```

```
182
183     # Put new_key back with updated freq
184     freqs[new_key] = freqs.get(new_key, 0) + old_freq
185
186     pairs_to_keys[pair] = set()
187     return changed_pairs
188
189
190 def write_merges(merges, outpath):
191     """Pickle the merges list to a binary file."""
192     os.makedirs(os.path.dirname(outpath), exist_ok=True)
193     with open(outpath, "wb") as f:
194         pickle.dump(merges, f)
195     print(f"Saved {len(merges)} merges to {outpath}")
196
197
198 def write_vocab(vocab, outpath):
199     """Pickle the vocab dict to a binary file."""
200     os.makedirs(os.path.dirname(outpath), exist_ok=True)
201     with open(outpath, "wb") as f:
202         pickle.dump(vocab, f)
203     print(f"Saved vocabulary with {len(vocab)} tokens to {outpath}")
204
205
206 def train_bpe(
207     input_path: str,
208     vocab_size: int,
209     special_tokens: list[str],
210     merges_outpath: str = None,
211     vocab_outpath: str = None,
212 ) -> tuple[dict[int, bytes], list[tuple[bytes, bytes]]]:
213     """
214     Trains byte-level BPE on a text file, returning:
215     - vocab: dict[int, bytes]
216     - merges: list of merged pairs
217     """
218     train_start_time = time.time()
219
220     # Initialize special tokens and single-byte tokens
221     initial_tokens = [tok.encode("UTF-8") for tok in special_tokens] + [bytes([i])
for i in range(256)]
222     vocab = {i: token for i, token in enumerate(initial_tokens)}
223     merges = []
224
225     print("Pre-tokenize: start")
226     start_time = time.time()
227     freqs = pre_tokenize(input_path, special_tokens)
228     print(f"Pre-tokenize: finished in {time.time() - start_time:.2f}s")
229
230     print("Initial pair frequencies: start")
```

```

231     start_time = time.time()
232     pair_freqs, pairs_to_keys = get_pair_freqs(freqs)
233
234     # Build a max-heap by pushing negative frequencies, ReverseLexOrderPair resolves
ties with max lexicographic order (hence the reverse)
235     pair_heap = []
236     for p, f in pair_freqs.items():
237         if f > 0:
238             heapq.heappush(pair_heap, (-f, ReverseLexOrderPair(p), p))
239
240     print(f"Initial pair frequencies: finished in {time.time() - start_time:.2f}s")
241
242     n_initial_tokens = len(initial_tokens)
243     n_merges = vocab_size - n_initial_tokens
244
245     print("Merge: start")
246     start_time = time.time()
247
248     for i in range(n_initial_tokens, n_initial_tokens + n_merges):
249         if not pair_heap:
250             break
251
252         # Pop until we find the top pair that still matches pair_freqs
253         while pair_heap:
254             neg_freq, _, top_pair = heapq.heappop(pair_heap)
255             freq = -neg_freq
256             if pair_freqs.get(top_pair, 0) == freq:
257                 pair = top_pair
258                 break
259             if top_pair in pair_freqs and pair_freqs[top_pair] > 0:
260                 heapq.heappush(pair_heap, (-pair_freqs[top_pair],
ReverseLexOrderPair(top_pair), top_pair))
261         else:
262             # If pair_heap is empty after the loop, we are done
263             break
264
265         if pair_freqs.get(pair, 0) <= 0:
266             break
267
268         # Add this new merge token to vocab and record the merge
269         vocab[i] = pair[0] + pair[1]
270         merges.append(pair)
271
272         # Merge in freqs, then update the heap for pairs changed by this merge
273         changed_pairs = merge(freqs, pair_freqs, pairs_to_keys, pair)
274         for cp in changed_pairs:
275             if cp in pair_freqs and pair_freqs[cp] > 0:
276                 heapq.heappush(pair_heap, (-pair_freqs[cp], ReverseLexOrderPair(cp),
cp))
277
278         # Print progress every 100 merges or at the last iteration

```

```
279         if ((i > n_initial_tokens) and ((i - n_initial_tokens + 1) % 100 == 0)) or (
280             i == n_initial_tokens + n_merges - 1
281         ):
282             print(
283                 f"{i - n_initial_tokens + 1}/{n_merges} merges completed (merge
runtime: {time.time() - start_time:.2f}s)"
284             )
285
286         print(f"Merges completed in {time.time() - start_time:.2f}s")
287         print(f"Training completed in {time.time() - train_start_time:.2f}s")
288
289         # Optionally save merges and vocab
290         if merges_outpath:
291             write_merges(merges, merges_outpath)
292         if vocab_outpath:
293             write_vocab(vocab, vocab_outpath)
294
295         return vocab, merges
296
297
298 if __name__ == "__main__":
299     (vocab, merges) = train_bpe(
300         input_path="./data/TinyStoriesV2-GPT4-valid.txt",
301         vocab_size=10000,
302         special_tokens=["<|endoftext|>"],
303         merges_outpath="./out/ts-valid-merges-2.txt",
304         vocab_outpath="./out/ts-valid-vocab-2.txt",
305     )
306
```

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snyder/cs336\_basics/tokenizer.py

```

1  from collections.abc import Iterable, Iterator
2  import regex as re
3  import pickle
4
5  from cs336_basics import train_bpe
6
7
8  class Tokenizer:
9      def __init__(
10         self, vocab: dict[int, bytes], merges: list[tuple[bytes, bytes]],
11         special_tokens: list[str] | None = None
12     ):
13         """
14         Constructs a tokenizer from a vocab, list of merges, and (optionally) list of
15         special tokens.
16         """
17         self.vocab = vocab
18         self.vocab_inv = {v: k for k, v in vocab.items()}
19         self.merges = merges
20         self.merges_dict = {merge: i for i, merge in enumerate(merges)}
21         self.encode_cache = {}
22         self.cache_hits = 0
23
24         self.pretokenize_pattern = re.compile(train_bpe.PAT)
25
26         # If there are extra special tokens than those used to construct the
27         # vocabulary initially
28         if special_tokens:
29             self.special_tokens = sorted(special_tokens, key=len, reverse=True)
30             self.special_pattern = "(" + "|".join(re.escape(k) for k in
31             self.special_tokens) + ")"
32
33             next_id = max(self.vocab.keys()) + 1
34             for token in special_tokens:
35                 token_bytes = token.encode("UTF-8")
36                 if token_bytes not in self.vocab_inv:
37                     self.vocab[next_id] = token_bytes
38                     self.vocab_inv[token_bytes] = next_id
39                     next_id += 1
40         else:
41             self.special_tokens = None
42             self.special_pattern = None
43
44     @classmethod
45     def from_files(cls, vocab_filepath: str, merges_filepath: str, special_tokens:
46     list[str] | None = None):
47         """

```

```

43     Constructs a Tokenizer from a serialized vocab, list of merges, and
    (optionally) list of special tokens.
44     """
45     with open(vocab_filepath, "rb") as f:
46         vocab = pickle.load(f)
47
48     with open(merges_filepath, "rb") as f:
49         merges = pickle.load(f)
50
51     # This calls the regular constructor, called as
    Tokenizer.from_files(vocab_filepath, merges_filepath, special_tokens)
52     return cls(vocab, merges, special_tokens)
53
54     def encode(self, text: str) -> list[int]:
55         """Encodes an input text into a sequence of token IDs, handling special
    tokens."""
56         if not self.special_tokens:
57             return self._encode_chunk(text)
58
59         # If we have special tokens, split on them, keeping delimiters
60         special_chunks = re.split(self.special_pattern, text)
61
62         ids = []
63         for part in special_chunks:
64             if part in self.special_tokens:
65                 # this is a special token
66                 ids.append(self.vocab_inv[part.encode("UTF-8")])
67             else:
68                 # this is ordinary text
69                 ids.extend(self._encode_chunk(part))
70         return ids
71
72     def _encode_chunk(self, text: str) -> list[int]:
73         """Encodes an input text chunk into a sequence of token IDs."""
74         pretokens = self._pretokenize(text)
75         pretoken_reprs: dict[str, list[bytes]] = {}
76
77         ids = []
78
79         # Merge each pretoken using the BPE rules, in ascending rank order
80         for p in pretokens:
81             # Check if we have already merged this subword previously
82             if p in self.encode_cache:
83                 ids.extend(self.encode_cache[p])
84                 self.cache_hits += 1
85             else:
86                 # Each character -> single bytes: e.g. "abc" -> [b'a', b'b', b'c']
87                 if p not in pretoken_reprs:
88                     match_bytes = list(bytes([b]) for b in p.encode("UTF-8"))
89                     pretoken_reprs[p] = match_bytes
90

```

```

91         merged = self._merge_subword(pretoken_reprs[p])
92         token_ids = [self.vocab_inv[subword] for subword in merged]
93         self.encode_cache[p] = token_ids
94         ids.extend(token_ids)
95
96     return ids
97
98     def _merge_subword(self, rep: list[bytes]) -> list[bytes]:
99         """
100         Given a list of subword units (bytes), repeatedly merges adjacent pairs
101         in ascending rank order until no more merges are found.
102         """
103         while True:
104             best_rank = float("inf")
105             best_idx = None
106
107             # Scan adjacent pairs in current rep, finding the earliest merge
108             for i in range(len(rep) - 1):
109                 pair = (rep[i], rep[i + 1])
110                 rank = self.merges_dict.get(pair)
111                 if rank is not None and rank < best_rank:
112                     best_rank = rank
113                     best_idx = i
114
115             # If no merges found, we're done
116             if best_idx is None:
117                 return rep
118
119             # Merge the best pair
120             merged = rep[best_idx] + rep[best_idx + 1] # Concatenate bytes
121             rep = rep[:best_idx] + [merged] + rep[best_idx + 2 :]
122
123     def encode_iterable(self, iterable: Iterable[str]) -> Iterator[int]:
124         """Yields token IDs lazily from an iterable of strings (e.g., a file
125         handle)."""
126         for text in iterable:
127             yield from self.encode(text)
128
129     def decode(self, ids: list[int]) -> str:
130         """Decodes a sequence of token IDs into text."""
131         text = b"".join(self.vocab[id] for id in ids)
132         return text.decode("UTF-8", errors="replace")
133
134     def _pretokenize(self, text: str) -> list[str]:
135         """Splits text into 'pretokens' and builds an initial byte representation for
136         each."""
137         pretokens: list[str] = []
138
139         for match in self.pretokenize_pattern.finditer(text):
140             match_str = match.group()
141             pretokens.append(match_str)

```

140

141

142

**return** pretokens






Figure 1: An overview of our Transformer language model.

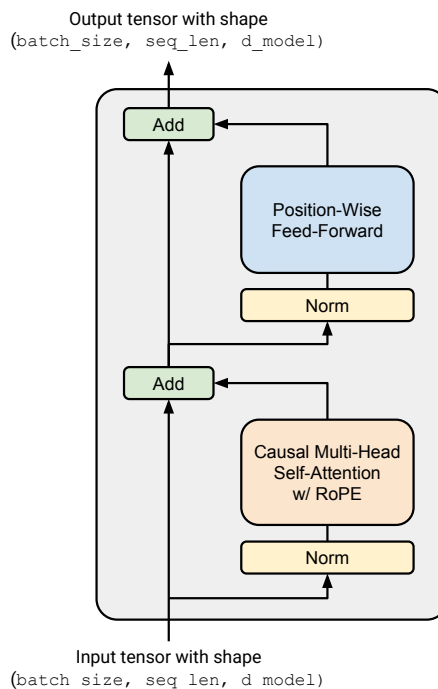


Figure 2: A pre-norm Transformer block.

### 3 Transformer Language Model Architecture

A language model takes as **input** a batched sequence of integer token IDs (i.e., `torch.Tensor` of shape `(batch_size, sequence_length)`), and **returns** a (batched) normalized probability distribution over the vocabulary (i.e., a PyTorch Tensor of shape `(batch_size, sequence_length, vocab_size)`), where the predicted distribution is over the next word for each input token. **When training the language model, we use these next-word predictions to calculate the cross-entropy loss between the actual next word and the predicted next word.** When generating text from the language model during inference, we take the predicted next-word distribution from the final time step (i.e., the last item in the sequence) to generate the next token in the sequence (e.g., by taking the token with the highest probability, sampling from the distribution, etc.), add the generated token to the input sequence, and repeat.

In this part of the assignment, you will build this Transformer language model from scratch. We will begin with a high-level description of the model before progressively detailing the individual components.

#### 3.1 Transformer LM

Given a sequence of token IDs, the Transformer language model uses an input embedding to convert token IDs to dense vectors, passes the embedded tokens through `num_layers` Transformer blocks, and then applies a learned linear projection (the “output embedding” or “LM head”) to produce the predicted next-token logits. See Figure 1 for a schematic representation.

##### 3.1.1 Token Embeddings

In the **very first step**, the Transformer **embeds** the (batched) sequence of token IDs into a sequence of vectors containing information on the token identity (red blocks in Figure 1).

More specifically, given a sequence of token IDs, the Transformer language model uses a token embedding layer to produce a sequence of vectors. Each embedding layer takes in a tensor of integers of shape `(batch_size, sequence_length)` and produces a sequence of vectors of shape `(batch_size, sequence_length, d_model)`.

### 3.1.2 Pre-norm Transformer Block

After embedding, the activations are processed by several identically structured neural net layers. A standard **decoder-only** Transformer language model consists of `num_layers` identical layers (commonly called Transformer “blocks”). Each Transformer block takes in an input of shape `(batch_size, sequence_length, d_model)` and returns an output of shape `(batch_size, sequence_length, d_model)`. **Each block aggregates information across the sequence (via self-attention) and non-linearly transforms it (via the feed-forward layers).**

## 3.2 Output Normalization and Embedding

After `num_layers` Transformer blocks, we will take the final activations and **turn them into a distribution over the vocabulary**.

We will implement the “pre-norm” Transformer block (detailed in §3.5), which additionally requires the use of layer normalization (detailed below) after the final Transformer block to ensure its outputs are properly scaled.

After this normalization, we will use a standard learned linear transformation to convert the output of the Transformer blocks into predicted next-token logits (see, e.g., Radford et al. [2018] equation 2).

## 3.3 Remark: Batching, Einsum and Efficient Computation

Throughout the Transformer, we will be performing the same computation applied to many batch-like inputs. Here are a few examples:

- **Elements of a batch:** we apply the same Transformer **forward** operation on each batch element.
- **Sequence length:** the “position-wise” operations like RMSNorm and feed-forward operate identically on each position of a sequence.
- **Attention heads:** the attention operation is batched across attention heads in a “multi-headed” attention operation.

It is useful to have an ergonomic way of performing such operations in a way that fully utilizes the GPU, and is easy to read and understand. Many PyTorch operations can take in excess “batch-like” dimensions at the start of a tensor and repeat/broadcast the operation across these dimensions efficiently.

For instance, say we are doing a position-wise, batched operation. We have a “data tensor”  $D$  of shape `(batch_size, sequence_length, d_model)`, and we would like to do a batched vector-matrix multiply against a matrix  $A$  of shape `(d_model, d_model)`. In this case,  $D @ A$  will do a batched matrix multiply, which is an efficient primitive in PyTorch, where the `(batch_size, sequence_length)` dimensions are batched over.

Because of this, it is helpful to assume that your functions may be given additional batch-like dimensions and to **keep those dimensions at the start of the PyTorch shape**. To organize tensors so they can be batched in this manner, they might need to be shaped using many steps of `view`, `reshape` and `transpose`. This can be a bit of a pain, and it often gets hard to read what the code is doing and what the shapes of your tensors are.

A more ergonomic option is to use *einsum notation* within `torch.einsum`, or rather use framework agnostic libraries like `einops` or `einx`. The two key ops are `einsum`, which can do tensor contractions with arbitrary dimensions of input tensors, and `rearrange`, which can reorder, concatenate, and split arbitrary

dimensions. It turns out almost all operations in machine learning are some combination of dimension juggling and tensor contraction with the occasional (usually pointwise) nonlinear function. This means that a lot of your code can be more readable and flexible when using einsum notation.

We **strongly** recommend learning and using einsum notation for the class. Students who have not been exposed to einsum notation before should use `einops` (docs [here](#)), and students who are already comfortable with `einops` should learn the more general `einx` ([here](#)).<sup>4</sup> Both packages are already installed in the environment we've supplied.

Here we give some examples of how einsum notation can be used. These are a supplement to the documentation for `einops`, which you should read first.

#### Example (einstein\_example1): Batched matrix multiplication with `einops.einsum`

```
import torch
from einops import rearrange, einsum

## Basic implementation
Y = D @ A.T
# Hard to tell the input and output shapes and what they mean.
# What shapes can D and A have, and do any of these have unexpected behavior?

## Einsum is self-documenting and robust
#           D           A           ->           Y
Y = einsum(D, A, "batch sequence d_in, d_out d_in -> batch sequence d_out")

## Or, a batched version where D can have any leading dimensions but A is constrained.
Y = einsum(D, A, "... d_in, d_out d_in -> ... d_out")
```

#### Example (einstein\_example2): Broadcasted operations with `einops.rearrange`

We have a batch of images, and for each image we want to generate 10 dimmed versions based on some scaling factor:

```
images = torch.randn(64, 128, 128, 3)  # (batch, height, width, channel)
dim_by = torch.linspace(start=0.0, end=1.0, steps=10)

## Reshape and multiply
dim_value = rearrange(dim_by, "dim_value -> 1 dim_value 1 1 1")
images_rearr = rearrange(images, "b height width channel -> b 1 height width channel")
dimmed_images = images_rearr * dim_value

## Or in one go:
dimmed_images = einsum(
    images, dim_by,
    "batch height width channel, dim_value -> batch dim_value height width channel"
)
```

<sup>4</sup>It's worth noting that while `einops` has a great amount of support, `einx` is not as battle-tested. You should feel free to fall back to using `einops` with some more plain PyTorch if you find any limitations or bugs in `einx`.

### Example (einstein\_example3): Pixel mixing with einops.rearrange

Suppose we have a batch of images represented as a tensor of shape (batch, height, width, channel), and we want to perform a linear transformation across all pixels of the image, but this transformation should happen independently for each channel. Our linear transformation is represented as a matrix  $B$  of shape (height  $\times$  width, height  $\times$  width).

```
channels_last = torch.randn(64, 32, 32, 3) # (batch, height, width, channel)
B = torch.randn(32*32, 32*32)

## Rearrange an image tensor for mixing across all pixels
channels_last_flat = channels_last.view(
    -1, channels_last.size(1) * channels_last.size(2), channels_last.size(3)
)
channels_first_flat = channels_last_flat.transpose(1, 2)

channels_first_flat_transformed = channels_first_flat @ B.T

channels_last_flat_transformed = channels_first_flat_transformed.transpose(1, 2)

channels_last_transformed = channels_last_flat_transformed.view(*channels_last.shape)
```

Instead, using einops:

```
height = width = 32
## Rearrange replaces clunky torch view + transpose
channels_first = rearrange(
    channels_last,
    "batch height width channel -> batch channel (height width)"
)
channels_first_transformed = einsum(
    channels_first, B,
    "batch channel pixel_in, pixel_out pixel_in -> batch channel pixel_out"
)
channels_last_transformed = rearrange(
    channels_first_transformed,
    "batch channel (height width) -> batch height width channel",
    height=height, width=width
)
```

Or, if you're feeling crazy: all in one go using einx.dot (einx equivalent of einops.einsum)

```
height = width = 32
channels_last_transformed = einx.dot(
    "batch row_in col_in channel, (row_out col_out) (row_in col_in)"
    "-> batch row_out col_out channel",
    channels_last, B,
    col_in=width, col_out=width
)
```

The first implementation here could be improved by placing comments before and after to indicate

what the input and output shapes are, but this is clunky and susceptible to bugs. With einsum notation, documentation *is* implementation!

Einsum notation can handle arbitrary input batching dimensions, **but also has the key benefit of being self-documenting**. It's much clearer what the relevant shapes of your input and output tensors are in code that uses einsum notation. For the remaining tensors, you can consider using Tensor type hints, for instance using the **jaxtyping** library (not specific to Jax).

We will talk more about the performance implications of using einsum notation in assignment 2, but for now know that they're almost always better than the alternative!

### 3.3.1 Mathematical Notation and Memory Ordering

Many machine learning papers use row vectors in their notation, which result in representations that mesh well with the row-major memory ordering used by default in NumPy and PyTorch. With row vectors, a linear transformation looks like

$$y = xW^\top, \tag{1}$$

for row-major  $W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$  and row-vector  $x \in \mathbb{R}^{1 \times d_{\text{in}}}$ .

In linear algebra it's generally more common to use column vectors, where linear transformations look like

$$y = Wx, \tag{2}$$

given a row-major  $W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$  and column-vector  $x \in \mathbb{R}^{d_{\text{in}}}$ . **We will use column vectors** for mathematical notation in this assignment, as it is generally easier to follow the math this way. **You should keep in mind that if you want to use plain matrix multiplication notation, you will have to apply matrices using the row vector convention, since PyTorch uses row-major memory ordering.** If you use einsum for your matrix operations, this should be a non-issue.

## 3.4 Basic Building Blocks: Linear and Embedding Modules

### 3.4.1 Parameter Initialization

Training neural networks effectively often requires careful initialization of the model parameters—bad initializations can lead to undesirable behavior such as vanishing or exploding gradients. Pre-norm transformers are unusually robust to initializations, but they can still have a significant impact on training speed and convergence. Since this assignment is already long, we will save the details for assignment 3, and instead give you some approximate initializations that should work well for most cases. For now, use:

- Linear weights:  $\mathcal{N}\left(\mu = 0, \sigma^2 = \frac{2}{d_{\text{in}} + d_{\text{out}}}\right)$  truncated at  $[-3\sigma, 3\sigma]$ .
- Embedding:  $\mathcal{N}(\mu = 0, \sigma^2 = 1)$  truncated at  $[-3, 3]$
- RMSNorm: 1

**You should use `torch.nn.init.trunc_normal_` to initialize the truncated normal weights.**

### 3.4.2 Linear Module

Linear layers are a fundamental building block of Transformers and neural nets in general. First, you will implement your own `Linear` class that inherits from `torch.nn.Module` and performs a linear transformation:

$$y = Wx. \tag{3}$$

Note that **we do not include a bias term, following most modern LLMs.**

### Problem (linear): Implementing the linear module (1 point)

**Deliverable:** Implement a `Linear` class that inherits from `torch.nn.Module` and performs a linear transformation. Your implementation should follow the interface of PyTorch's built-in `nn.Linear` module, except for not having a `bias` argument or parameter. We recommend the following interface:

```
def __init__(self, in_features, out_features, device=None, dtype=None) Construct a
    linear transformation module. This function should accept the following parameters:

    in_features: int final dimension of the input
    out_features: int final dimension of the output
    device: torch.device | None = None Device to store the parameters on
    dtype: torch.dtype | None = None Data type of the parameters

def forward(self, x: torch.Tensor) -> torch.Tensor Apply the linear transformation to the
    input.
```

Make sure to:

- subclass `nn.Module`
- call the superclass constructor
- construct and store your parameter as  $W$  (not  $W^T$ ) for memory ordering reasons, putting it in an `nn.Parameter`
- of course, don't use `nn.Linear` or `nn.functional.linear`

For initializations, use the settings from above along with `torch.nn.init.trunc_normal_` to initialize the weights.

To test your `Linear` module, implement the test adapter at `[adapters.run_linear]`. The adapter should load the given weights into your `Linear` module. You can use `Module.load_state_dict` for this purpose. Then, run `uv run pytest -k test_linear`.

### 3.4.3 Embedding Module

As discussed above, the first layer of the Transformer is an embedding layer that maps integer token IDs into a vector space of dimension `d_model`. We will implement a custom `Embedding` class that inherits from `torch.nn.Module` (so you should not use `nn.Embedding`). The `forward` method should select the embedding vector for each token ID by indexing into an embedding matrix of shape `(vocab_size, d_model)` using a `torch.LongTensor` of token IDs with shape `(batch_size, sequence_length)`.

### Problem (embedding): Implement the embedding module (1 point)

**Deliverable:** Implement the `Embedding` class that inherits from `torch.nn.Module` and performs an embedding lookup. Your implementation should follow the interface of PyTorch's built-in `nn.Embedding` module. We recommend the following interface:

```
def __init__(self, num_embeddings, embedding_dim, device=None, dtype=None) Construct
    an embedding module. This function should accept the following parameters:

    num_embeddings: int Size of the vocabulary
```

```

embedding_dim: int Dimension of the embedding vectors, i.e.,  $d_{\text{model}}$ 
device: torch.device | None = None Device to store the parameters on
dtype: torch.dtype | None = None Data type of the parameters

def forward(self, token_ids: torch.Tensor) -> torch.Tensor Lookup the embedding vectors
for the given token IDs.

```

Make sure to:

- subclass `nn.Module`
- call the superclass constructor
- initialize your embedding matrix as a `nn.Parameter`
- store the embedding matrix with the `d_model` being the final dimension
- of course, don't use `nn.Embedding` or `nn.functional.embedding`

Again, use the settings from above for initialization, and use `torch.nn.init.trunc_normal_` to initialize the weights.

To test your implementation, implement the test adapter at `[adapters.run_embedding]`. Then, run `uv run pytest -k test_embedding`.

## 3.5 Pre-Norm Transformer Block

Each Transformer block has two sub-layers: a multi-head self-attention mechanism and a position-wise feed-forward network (Vaswani et al., 2017, section 3.1).

In the original Transformer paper, the model uses a residual connection around each of the two sub-layers, followed by layer normalization. This architecture is commonly known as the “post-norm” Transformer, since layer normalization is applied to the sublayer output. However, a variety of work has found that moving layer normalization from the output of each sub-layer to the input of each sub-layer (with an additional layer normalization after the final Transformer block) improves Transformer training stability (Nguyen and Salazar, 2019; Xiong et al., 2020)—see Figure 2 for a visual representation of this “pre-norm” Transformer block. The output of each Transformer block sub-layer is then added to the sub-layer input via the residual connection (Vaswani et al., 2017, section 5.4). An intuition for pre-norm is that there is a clean “residual stream” without any normalization going from the input embeddings to the final output of the Transformer, which is purported to improve gradient flow. This pre-norm Transformer is now the standard used in language models today (e.g., GPT-3, LLaMA, PaLM, etc.), so we will implement this variant. We will walk through each of the components of a pre-norm Transformer block, implementing them in sequence.

### 3.5.1 Root Mean Square Layer Normalization

The original Transformer implementation of Vaswani et al. (2017) uses layer normalization (Ba et al., 2016) to normalize activations. Following Touvron et al. (2023), we will use root mean square layer normalization (RMSNorm; Zhang and Sennrich, 2019, equation 4) for layer normalization. Given a vector  $a \in \mathbb{R}^{d_{\text{model}}}$  of activations, RMSNorm will rescale each activation  $a_i$  as follows:

$$\text{RMSNorm}(a_i) = \frac{a_i}{\text{RMS}(a)} g_i, \quad (4)$$

where  $\text{RMS}(a) = \sqrt{\frac{1}{d_{\text{model}}} \sum_{i=1}^{d_{\text{model}}} a_i^2 + \varepsilon}$ . Here,  $g_i$  is a learnable “gain” parameter (there are `d_model` such parameters total), and  $\varepsilon$  is a hyperparameter that is often fixed at  $1e-5$ .

You should upcast your input to `torch.float32` to prevent overflow when you square the input. Overall, your `forward` method should look like:

```
in_dtype = x.dtype
x = x.to(torch.float32)

# Your code here performing RMSNorm
...
result = ...

# Return the result in the original dtype
return result.to(in_dtype)
```

#### Problem (rmsnorm): Root Mean Square Layer Normalization (1 point)

**Deliverable:** Implement `RMSNorm` as a `torch.nn.Module`. We recommend the following interface:

```
def __init__(self, d_model: int, eps: float = 1e-5, device=None, dtype=None)
    Construct the RMSNorm module. This function should accept the following parameters:

    d_model: int Hidden dimension of the model
    eps: float = 1e-5 Epsilon value for numerical stability
    device: torch.device | None = None Device to store the parameters on
    dtype: torch.dtype | None = None Data type of the parameters

def forward(self, x: torch.Tensor) -> torch.Tensor Process an input tensor of shape
(batch_size, sequence_length, d_model) and return a tensor of the same shape.
```

**Note:** Remember to upcast your input to `torch.float32` before performing the normalization (and later downcast to the original dtype), as described above.  
To test your implementation, implement the test adapter at `[adapters.run_rmsnorm]`. Then, run `uv run pytest -k test_rmsnorm`.



### 3.5.2 Position-Wise Feed-Forward Network




Figure 3: Comparing the SiLU (aka Swish) and ReLU activation functions.

In the original Transformer paper (section 3.3 of Vaswani et al. [2017]), the Transformer feed-forward network consists of two linear transformations with a ReLU activation ( $\text{ReLU}(x) = \max(0, x)$ ) between them. The dimensionality of the inner feed-forward layer is typically 4x the input dimensionality.

However, modern language models tend to incorporate two main changes compared to this original design: they use another activation function and employ a gating mechanism. Specifically, we will implement the “SwiGLU” activation function adopted in LLMs like Llama 3 [Grattafiori et al., 2024] and Qwen 2.5 [Yang et al., 2024], which combines the SiLU (often called Swish) activation with a gating mechanism called a Gated Linear Unit (GLU). We will also omit the bias terms sometimes used in linear layers, following most modern LLMs since PaLM [Chowdhery et al., 2022] and LLaMA [Touvron et al., 2023].

The SiLU or Swish activation function [Hendrycks and Gimpel, 2016, Elfwing et al., 2017] is defined as follows:

$$\text{SiLU}(x) = x \cdot \sigma(x) = \frac{x}{1 + e^{-x}} \quad (5)$$

As can be seen in Figure 3, the SiLU activation function is similar to the ReLU activation function, but is smooth at zero.

Gated Linear Units (GLUs) were originally defined by Dauphin et al. [2017] as the element-wise product of a linear transformation passed through a sigmoid function and another linear transformation:

$$\text{GLU}(x, W_1, W_2) = \sigma(W_1 x) \odot W_2 x, \quad (6)$$

where  $\odot$  represents element-wise multiplication. Gated Linear Units are suggested to “reduce the vanishing gradient problem for deep architectures by providing a linear path for the gradients while retaining non-linear capabilities.”

Putting the SiLU/Swish and GLU together, we get the SwiGLU, which we will use for our feed-forward networks:

$$\text{FFN}(x) = \text{SwiGLU}(x, W_1, W_2, W_3) = W_2(\text{SiLU}(W_1 x) \odot W_3 x), \quad (7)$$

where  $x \in \mathbb{R}^{d_{\text{model}}}$ ,  $W_1, W_3 \in \mathbb{R}^{d_{\text{ff}} \times d_{\text{model}}}$ ,  $W_2 \in \mathbb{R}^{d_{\text{model}} \times d_{\text{ff}}}$ , and canonically,  $d_{\text{ff}} = \frac{8}{3}d_{\text{model}}$ .

Shazeer [2020] first proposed combining the SiLU/Swish activation with GLUs and conducted experiments showing that SwiGLU outperforms baselines like ReLU and SiLU (without gating) on language modeling tasks. Later in the assignment, you will compare SwiGLU and SiLU. Though we've mentioned some heuristic arguments for these components (and the papers provide more supporting evidence), it's good to keep an empirical perspective: a now famous quote from Shazeer's paper is

We offer no explanation as to why these architectures seem to work; we attribute their success, as all else, to divine benevolence.

**Problem (positionwise\_feedforward): Implement the position-wise feed-forward network (2 points)**

**Deliverable:** Implement the SwiGLU feed-forward network, composed of a SiLU activation function and a GLU.

**Note:** in this particular case, you should feel free to use `torch.sigmoid` in your implementation for numerical stability.

You should set  $d_{\text{ff}}$  to approximately  $\frac{8}{3} \times d_{\text{model}}$  in your implementation, while ensuring that the dimensionality of the inner feed-forward layer is a multiple of 64 to make good use of your hardware. To test your implementation against our provided tests, you will need to implement the test adapter at `[adapters.run_swiglu]`. Then, run `uv run pytest -k test_swiglu` to test your implementation.

### 3.5.3 Relative Positional Embeddings

To inject positional information into the model, we will implement Rotary Position Embeddings [Su et al., 2021], often called RoPE. For a given query token  $q^{(i)} = W_q x^{(i)} \in \mathbb{R}^d$  at token position  $i$ , we will apply a pairwise rotation matrix  $R^i$ , giving us  $q'^{(i)} = R^i q^{(i)} = R^i W_q x^{(i)}$ . Here,  $R^i$  will rotate pairs of embedding elements  $q_{2k-1:2k}^{(i)}$  as 2d vectors by the angle  $\theta_{i,k} = \frac{i}{\Theta^{2k/d}}$  for  $k \in \{1, \dots, d/2\}$  and some constant  $\Theta$ . Thus, we can consider  $R^i$  to be a block-diagonal matrix of size  $d \times d$ , with blocks  $R_k^i$  for  $k \in \{1, \dots, d/2\}$ , with

$$R_k^i = \begin{bmatrix} \cos(\theta_{i,k}) & -\sin(\theta_{i,k}) \\ \sin(\theta_{i,k}) & \cos(\theta_{i,k}) \end{bmatrix}. \quad (8)$$

Thus we get the full rotation matrix

Should k start with 0 or with 1?

It probably doesn't matter...

d = d\_k

$$R^i = \begin{bmatrix} R_1^i & 0 & 0 & \dots & 0 \\ 0 & R_2^i & 0 & \dots & 0 \\ 0 & 0 & R_3^i & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & R_{d/2}^i \end{bmatrix}, \quad (9)$$

where 0s represent  $2 \times 2$  zero matrices. While one could construct the full  $d \times d$  matrix, a good solution should use the properties of this matrix to implement the transformation more efficiently. Since we only care about the relative rotation of tokens within a given sequence, we can reuse the values we compute for  $\cos(\theta_{i,k})$  and  $\sin(\theta_{i,k})$  across layers, and different batches. If you would like to optimize it, you may use a single RoPE module referenced by all layers, and it can have a 2d pre-computed buffer of sin and cos values created during init with `self.register_buffer(persistent=False)`, instead of a `nn.Parameter` (because we do not want to learn these fixed cosine and sine values). The exact same rotation process we did for our  $q^{(i)}$  is then done for  $k^{(j)}$ , rotating by the corresponding  $R^j$ . Notice that this layer has no learnable parameters.

### Problem (rope): Implement RoPE (2 points)

**Deliverable:** Implement a class `RotaryPositionalEmbedding` that applies RoPE to the input tensor.

The following interface is recommended:

```
def __init__(self, theta: float, d_k: int, max_seq_len: int, device=None) Construct the
    RoPE module and create buffers if needed.

    theta: float  $\Theta$  value for the RoPE
    d_k: int dimension of query and key vectors
    max_seq_len: int Maximum sequence length that will be inputted
    device: torch.device | None = None Device to store the buffer on

def forward(self, x: torch.Tensor, token_positions: torch.Tensor) -> torch.Tensor Process
    an input tensor of shape  $(\dots, \text{seq\_len}, d_k)$  and return a tensor of the same shape. Note
    that you should tolerate  $x$  with an arbitrary number of batch dimensions. You should assume
    that the token positions are a tensor of shape  $(\dots, \text{seq\_len})$  specifying the token positions of
     $x$  along the sequence dimension.

    You should use the token positions to slice your (possibly precomputed) cos and sin tensors along
    the sequence dimension.

    To test your implementation, complete [adapters.run_rope] and make sure it passes uv run
    pytest -k test_rope.
```

### 3.5.4 Scaled Dot-Product Attention

We will now implement scaled dot-product attention as described in [\[Vaswani et al. 2017\]](#) (section 3.2.1). As a preliminary step, the definition of the Attention operation will make use of softmax, an operation that takes an unnormalized vector of scores and turns it into a normalized distribution:

$$\text{softmax}(v)_i = \frac{\exp(v_i)}{\sum_{j=1}^n \exp(v_j)}. \quad (10)$$

Note that  $\exp(v_i)$  can become `inf` for large values (then, `inf/inf = NaN`). We can avoid this by noticing that the softmax operation is invariant to adding any constant  $c$  to all inputs. We can leverage this property for numerical stability—typically, we will subtract the largest entry of  $o_i$  from all elements of  $o_i$ , making the new largest entry 0. You will now implement softmax, using this trick for numerical stability.

### Problem (softmax): Implement softmax (1 point)

**Deliverable:** Write a function to apply the softmax operation on a tensor. Your function should take two parameters: a tensor and a *dimension*  $i$ , and apply softmax to the  $i$ -th dimension of the input tensor. The output tensor should have the same shape as the input tensor, but its  $i$ -th dimension will now have a normalized probability distribution. Use the trick of subtracting the maximum value in the  $i$ -th dimension from all elements of the  $i$ -th dimension to avoid numerical stability issues.

To test your implementation, complete `[adapters.run_softmax]` and make sure it passes `uv run pytest -k test_softmax_matches_pytorch`.

We can now define the Attention operation mathematically as follows:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q^\top K}{\sqrt{d_k}}\right) V \quad (11)$$

where  $Q \in \mathbb{R}^{n \times d_k}$ ,  $K \in \mathbb{R}^{m \times d_k}$ , and  $V \in \mathbb{R}^{m \times d_v}$ . Here,  $Q$ ,  $K$  and  $V$  are all inputs to this operation—note that these are not the learnable parameters. If you’re wondering why this isn’t  $QK^\top$ , see 3.3.1.

**Masking:** It is sometimes convenient to *mask* the output of an attention operation. A mask should have the shape  $M \in \{\text{True}, \text{False}\}^{n \times m}$ , and each row  $i$  of this boolean matrix indicates which keys the query  $i$  should attend to. Canonically (and slightly confusingly), a value of **True** at position  $(i, j)$  indicates that the query  $i$  *does* attend to the key  $j$ , and a value of **False** indicates that the query *does not* attend to the key. In other words, “information flows” at  $(i, j)$  pairs with value **True**. For example, consider a  $1 \times 3$  mask matrix with entries `[[True, True, False]]`. The single query vector attends only to the first two keys.

Computationally, it will be much more efficient to use masking than to compute attention on subsequences, and we can do this by taking the pre-softmax values  $\left(\frac{Q^\top K}{\sqrt{d_k}}\right)$  and adding a  $-\infty$  in any entry of the mask matrix that is False.

**Problem (scaled\_dot\_product\_attention): Implement scaled dot-product attention (5 points)**

**Deliverable:** Implement the scaled dot-product attention function. Your implementation should handle keys and queries of shape `(batch_size, ..., seq_len, d_k)` and values of shape `(batch_size, ..., seq_len, d_v)`, where `...` represents any number of other batch-like dimensions (if provided). The implementation should return an output with the shape `(batch_size, ..., d_v)`. See section 3.3 for a discussion on batch-like dimensions.

Your implementation should also support an optional user-provided boolean mask of shape `(seq_len, seq_len)`. The attention probabilities of positions with a mask value of **True** should collectively sum to 1, and the attention probabilities of positions with a mask value of **False** should be zero.

To test your implementation against our provided tests, you will need to implement the test adapter at `[adapters.run_scaled_dot_product_attention]`.

`uv run pytest -k test_scaled_dot_product_attention` tests your implementation on third-order input tensors, while `uv run pytest -k test_4d_scaled_dot_product_attention` tests your implementation on fourth-order input tensors.

### 3.5.5 Causal Multi-Head Self-Attention

We will implement multi-head self-attention as described in section 3.2.2 of Vaswani et al. [2017]. Recall that, mathematically, the operation of applying multi-head attention is defined as follows:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) \quad (12)$$

$$\text{for head}_i = \text{Attention}(Q_i, K_i, V_i) \quad (13)$$

with  $Q_i$ ,  $K_i$ ,  $V_i$  being slice number  $i \in \{1, \dots, h\}$  of size  $d_k$  or  $d_v$  of the embedding dimension for  $Q$ ,  $K$ , and  $V$  respectively. With Attention being the scaled dot-product attention operation defined in §3.5.4. From this we can form the multi-head *self*-attention operation:

$$\text{MultiHeadSelfAttention}(x) = W_O \text{MultiHead}(W_Q x, W_K x, W_V x) \quad (14)$$

Here, the learnable parameters are  $W_Q \in \mathbb{R}^{hd_k \times d_{\text{model}}}$ ,  $W_K \in \mathbb{R}^{hd_k \times d_{\text{model}}}$ ,  $W_V \in \mathbb{R}^{hd_v \times d_{\text{model}}}$ , and  $W_O \in \mathbb{R}^{d_{\text{model}} \times hd_v}$ . Since the  $Q$ s,  $K$ , and  $V$ s are sliced in the multi-head attention operation, we can think of  $W_Q$ ,  $W_K$  and  $W_V$  as being separated for each head along the output dimension. When you have this working, you should be computing the key, value, and query projections in a total of three matrix multiplies.<sup>5</sup>

<sup>5</sup>As a stretch goal, try combining the key, query, and value projections into a single weight matrix so you only need a single matrix multiply.

**Causal masking.** Your implementation should prevent the model from attending to future tokens in the sequence. In other words, if the model is given a token sequence  $t_1, \dots, t_n$ , and we want to calculate the next-word predictions for the prefix  $t_1, \dots, t_i$  (where  $i < n$ ), the model should *not* be able to access (attend to) the token representations at positions  $t_{i+1}, \dots, t_n$  since it will not have access to these tokens when generating text during inference (and these future tokens leak information about the identity of the true next word, trivializing the language modeling pre-training objective). For an input token sequence  $t_1, \dots, t_n$  we can naively prevent access to future tokens by running multi-head self-attention  $n$  times (for the  $n$  unique prefixes in the sequence). Instead, we'll use causal attention masking, which allows token  $i$  to attend to all positions  $j \leq i$  in the sequence. You can use `torch.triu` or a broadcasted index comparison to construct this `mask`, and you should take advantage of the fact that your scaled dot-product attention implementation from §3.5.4 already supports attention masking.

**Applying RoPE.** RoPE should be applied to the query and key vectors, but not the value vectors. Also, the head dimension should be handled as a batch dimension, because in multi-head attention, attention is being applied independently for each head. This means that precisely the same RoPE rotation should be applied to the query and key vectors for each head.

**Problem (multihead\_self\_attention): Implement causal multi-head self-attention (5 points)**

**Deliverable:** Implement causal multi-head self-attention as a `torch.nn.Module`. Your implementation should accept (at least) the following parameters:

**d\_model:** `int` Dimensionality of the Transformer block inputs.

**num\_heads:** `int` Number of heads to use in multi-head self-attention.

Following Vaswani et al. [2017], set  $d_k = d_v = d_{\text{model}}/h$ . To test your implementation against our provided tests, implement the test adapter at `[adapters.run_multihead_self_attention]`. Then, run `uv run pytest -k test_multihead_self_attention` to test your implementation.

### 3.6 The Full Transformer LM

Let's begin by assembling the Transformer block (it will be helpful to refer back to Figure 2). A Transformer block contains two 'sublayers', one for the multihead self attention, and another for the feed-forward network. In each sublayer, we first perform RMSNorm, then the main operation (MHA/FF), finally adding in the residual connection.

To be concrete, the first half (the first 'sub-layer') of the Transformer block should be implementing the following set of updates to produce an output  $y$  from an input  $x$ ,

$$y = x + \text{MultiHeadSelfAttention}(\text{RMSNorm}(x)). \quad (15)$$

**Problem (transformer\_block): Implement the Transformer block (3 points)**

Implement the pre-norm Transformer block as described in §3.5 and illustrated in Figure 2. Your Transformer block should accept (at least) the following parameters.

**d\_model:** `int` Dimensionality of the Transformer block inputs.

**num\_heads:** `int` Number of heads to use in multi-head self-attention.

**d\_ff:** `int` Dimensionality of the position-wise feed-forward inner layer.

To test your implementation, implement the adapter `[adapters.run_transformer_block]`. Then run `uv run pytest -k test_transformer_block` to test your implementation.

**Deliverable:** Transformer block code that passes the provided tests.

Now we put the blocks together, following the high level diagram in Figure 1. Follow our description of the embedding in Section 3.1.1, feed this into `num_layers` Transformer blocks, and then pass that into the three output layers to obtain a distribution over the vocabulary.

### Problem (transformer\_lm): Implementing the Transformer LM (3 points)

Time to put it all together! Implement the Transformer language model as described in §3.1 and illustrated in Figure 1. At minimum, your implementation should accept all the aforementioned construction parameters for the Transformer block, as well as these additional parameters:

**vocab\_size:** `int` The size of the vocabulary, necessary for determining the dimensionality of the token embedding matrix.

**context\_length:** `int` The maximum context length, necessary for determining the dimensionality of the position embedding matrix.

**num\_layers:** `int` The number of Transformer blocks to use.

To test your implementation against our provided tests, you will first need to implement the test adapter at `[adapters.run_transformer_lm]`. Then, run `uv run pytest -k test_transformer_lm` to test your implementation.

**Deliverable:** A Transformer LM module that passes the above tests.

**Resource accounting.** It is useful to be able to understand how the various parts of the Transformer consume compute and memory. We will go through the steps to do some basic “FLOPs accounting.” The vast majority of FLOPs in a Transformer are matrix multiplies, so our core approach is simple:

1. Write down all the matrix multiplies in a Transformer forward pass.
2. Convert each matrix multiply into FLOPs required.

For this second step, the following fact will be useful:

**Rule:** Given  $A \in \mathbb{R}^{m \times n}$  and  $B \in \mathbb{R}^{n \times p}$ , the matrix-matrix product  $AB$  requires  $2mnp$  FLOPs.

To see this, note that  $(AB)[i, j] = A[i, :] \cdot B[:, j]$ , and that this dot product requires  $n$  additions and  $n$  multiplications ( $2n$  FLOPs). Then, since the matrix-matrix product  $AB$  has  $m \times p$  entries, the total number of FLOPs is  $(2n)(mp) = 2mnp$ .

Now, before you do the next problem, it can be helpful to go through each component of your Transformer block and Transformer LM, and list out all the matrix multiplies and their associated FLOPs costs.

### Problem (transformer\_accounting): Transformer LM resource accounting (5 points)

(a) Consider GPT-2 XL, which has the following configuration:

<b>vocab_size</b> : 50,257	$M_{\text{emb}} = d_v \times d_{\text{model}} = 50,257 \times 1600$
	$M_q = d_{\text{model}} \times d_{\text{quer}} (\text{tot}) = d_{\text{model}} \times d_{\text{model}} = 1600 \times 1600$
<b>context_length</b> : 1,024	$M_k = d_{\text{model}} \times d_{\text{model}} = 1600 \times 1600$
	$M_v = d_{\text{model}} \times d_{\text{model}} = 1600 \times 1600$
<b>num_layers</b> : 48	$M_{W1} = d_{\text{model}} \times d_{\text{ff}} = 1600 \times 6400$
	$M_{W2} = d_{\text{model}} \times d_{\text{ff}} = 1600 \times 6400$
<b>d_model</b> : 1,600	$M_{W3} = d_{\text{model}} \times d_{\text{ff}} = 1600 \times 6400$
	$M_{\text{rot}} = 2 \times 2 \times \text{context\_length} / 2 = 2 \times \text{context\_length} = 2048$ , Not Trainable!

**num\_heads** : 25

**d\_ff** : 6,400

Suppose we constructed our model using this configuration. How many trainable parameters would our model have? Assuming each parameter is represented using single-precision floating point, how much memory is required to just load this model?

**Deliverable:** A one-to-two sentence response.

- (b) Identify the matrix multiplies required to complete a forward pass of our GPT-2 XL-shaped model. How many FLOPs do these matrix multiplies require in total? Assume that our input sequence has **context\_length** tokens.

**Deliverable:** A list of matrix multiplies (with descriptions), and the total number of FLOPs required.

- (c) Based on your analysis above, which parts of the model require the most FLOPs?

**Deliverable:** A one-to-two sentence response.

- (d) Repeat your analysis with GPT-2 small (12 layers, 768 **d\_model**, 12 heads), GPT-2 medium (24 layers, 1024 **d\_model**, 16 heads), and GPT-2 large (36 layers, 1280 **d\_model**, 20 heads). As the model size increases, which parts of the Transformer LM take up proportionally more or less of the total FLOPs?

**Deliverable:** For each model, provide a breakdown of model components and its associated FLOPs (as a proportion of the total FLOPs required for a forward pass). In addition, provide a one-to-two sentence description of how varying the model size changes the proportional FLOPs of each component.

- (e) Take GPT-2 XL and increase the context length to 16,384. How does the total FLOPs for one forward pass change? How do the relative contribution of FLOPs of the model components change?

**Deliverable:** A one-to-two sentence response.

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snider/cs336\_basics/model.py

```

1 import torch
2 import math
3 from einops import einsum, rearrange, reduce
4
5
6 def softmax(x: torch.Tensor, dim: int) -> torch.Tensor:
7     x_max = x.max(dim=dim, keepdim=True).values
8     x_exp = torch.exp(x - x_max)
9     return x_exp / x_exp.sum(dim=dim, keepdim=True)
10
11
12 def scaled_dot_product_attention(Q: torch.Tensor, K: torch.Tensor, V: torch.Tensor,
13 mask: torch.Tensor):
14     d_k = Q.shape[-1]
15
16     attention_scores = einsum(Q, K, "... seq_q d, ... seq_k d -> ... seq_q seq_k")
17     attention_scores = attention_scores / math.sqrt(d_k)
18     attention_scores = torch.where(mask, attention_scores, float("-inf"))
19
20     attention_weights = softmax(attention_scores, dim=-1)
21     output = einsum(attention_weights, V, "... seq_q seq_k, ... seq_k d -> ... seq_q
22 d")
23
24     return output
25
26
27 class Linear(torch.nn.Module):
28     def __init__(
29         self, in_features: int, out_features: int, device: torch.device | None =
30 None, dtype: torch.dtype | None = None
31     ):
32         super().__init__()
33
34         mean = 0
35         std = math.sqrt(2 / (out_features + in_features))
36         lower = -3 * std
37         upper = 3 * std
38
39         w = torch.empty((out_features, in_features), device=device, dtype=dtype)
40         torch.nn.init.trunc_normal_(w, mean=mean, std=std, a=lower, b=upper)
41
42         self.weight = torch.nn.Parameter(w)
43
44
45     def forward(self, x: torch.Tensor) -> torch.Tensor:
46         return einsum(self.weight, x, "d_out d_in, ... d_in -> ... d_out")
47
48
49 class RotaryPositionalEmbedding(torch.nn.Module):

```



```

46     def __init__(
47         self,
48         theta: float,
49         d_k: int,
50         max_seq_len: int,
51         device: torch.device | None = None,
52         dtype: torch.dtype | None = None,
53     ):
54         super().__init__()
55
56         positions = torch.arange(max_seq_len, device=device).unsqueeze(1)
57         freqs = torch.arange(0, d_k, 2, device=device) / d_k # should this start with
a 1?
58         inv_freq = 1.0 / (theta**freqs)
59         angles = positions * inv_freq
60
61         self.register_buffer("cos", angles.cos().to(dtype), persistent=False)
62         self.register_buffer("sin", angles.sin().to(dtype), persistent=False)
63
64     def forward(self, x: torch.Tensor, token_positions: torch.Tensor) ->
torch.Tensor:
65         cos_pos = self.cos[token_positions]
66         sin_pos = self.sin[token_positions]
67
68         x_even = x[..., 0::2]
69         x_odd = x[..., 1::2]
70
71         x_rot_even = x_even * cos_pos - x_odd * sin_pos
72         x_rot_odd = x_even * sin_pos + x_odd * cos_pos
73
74         x_rot = rearrange([x_rot_even, x_rot_odd], "two ... -> ... two")
75         x_out = rearrange(x_rot, "... d1 d2 -> ... (d1 d2)")
76
77         return x_out
78
79
80 class CausalMultiHeadSelfAttention(torch.nn.Module):
81     def __init__(self, d_model: int, num_heads: int, device=None, dtype=None,
**kwargs):
82         super().__init__()
83
84         self.wqkv = Linear(d_model, 3 * d_model, device, dtype)
85         self.output_proj = Linear(d_model, d_model, device, dtype)
86
87         self.num_heads = num_heads
88         self.d_model = d_model
89         self.d_head = d_model // num_heads
90
91     def forward(
92         self,
93         x: torch.Tensor,

```

```

94     rope: RotaryPositionalEmbedding | None = None,
95     token_positions: torch.Tensor | None = None,
96 ) -> torch.Tensor:
97     batch_size, seq_len, _ = x.shape
98     qkv = self.wqkv(x)
99
100     # Split into separate q, k, v tensors
101     q, k, v = qkv.split(self.d_model, dim=2)
102
103     # Reshape from (batch, seq_len, dim) to (batch, heads, seq_len, head_dim)
104     q = rearrange(q, "b s (h d) -> b h s d", h=self.num_heads)
105     k = rearrange(k, "b s (h d) -> b h s d", h=self.num_heads)
106     v = rearrange(v, "b s (h d) -> b h s d", h=self.num_heads)
107
108     if rope is not None:
109         if token_positions is None:
110             token_positions = torch.arange(seq_len, device=x.device)
111             q = rope(q, token_positions)
112             k = rope(k, token_positions)
113
114     # Create causal mask for self-attention
115     mask = ~torch.triu(torch.ones((seq_len, seq_len), device=x.device,
dtype=torch.bool), diagonal=1)
116
117     y = scaled_dot_product_attention(q, k, v, mask)
118     y = rearrange(y, "b h s d -> b s (h d)")
119     return self.output_proj(y)
120
121
122 class RMSNorm(torch.nn.Module):
123     def __init__(self, d_model: int, eps: float = 1e-5, device=None, dtype=None):
124         super().__init__()
125
126         self.eps = eps
127         self.weight = torch.nn.Parameter(torch.ones(d_model, device=device,
dtype=dtype))
128
129     def forward(self, x: torch.Tensor) -> torch.Tensor:
130         in_dtype = x.dtype
131         x = x.to(torch.float32)
132
133         rms = torch.sqrt(reduce(x**2, "... d -> ... 1", "mean") + self.eps)
134         result = x * self.weight / rms
135
136         return result.to(in_dtype)
137
138
139 def silu_activation(x: torch.Tensor) -> torch.Tensor:
140     return x * torch.sigmoid(x)
141
142

```

```
143 class SwiGLU(torch.nn.Module):
144     def __init__(self, d_model: int, d_ff: int, device: torch.device | None = None,
dtype: torch.dtype | None = None):
145         super().__init__()
146
147         self.w1 = Linear(d_model, d_ff, device, dtype)
148         self.w2 = Linear(d_ff, d_model, device, dtype)
149         self.w3 = Linear(d_model, d_ff, device, dtype)
150
151     def forward(self, x: torch.Tensor) -> torch.Tensor:
152         a1 = self.w1(x)
153         silu = silu_activation(a1)
154         return self.w2(silu * self.w3(x))
155
156
157 class SiLU(torch.nn.Module):
158     def __init__(self, d_model: int, d_ff: int, device: torch.device | None = None,
dtype: torch.dtype | None = None):
159         super().__init__()
160
161         self.w1 = Linear(d_model, d_ff, device, dtype)
162         self.w2 = Linear(d_ff, d_model, device, dtype)
163
164     def forward(self, x: torch.Tensor) -> torch.Tensor:
165         a1 = self.w1(x)
166         silu = silu_activation(a1)
167         return self.w2(silu)
168
169
170 class Block(torch.nn.Module):
171     def __init__(
172         self,
173         d_model: int,
174         num_heads: int,
175         d_ff: int,
176         rope: RotaryPositionalEmbedding | None = None,
177         device=None,
178         dtype=None,
179         **kwargs,
180     ):
181         super().__init__()
182
183         self.rope = rope
184
185         self.ln1 = RMSNorm(d_model, device=device, dtype=dtype)
186         self.attn = CausalMultiHeadSelfAttention(d_model, num_heads, device, dtype,
**kwargs)
187
188         self.ln2 = RMSNorm(d_model, device=device, dtype=dtype)
189
190         ffn_type = kwargs.get("ffn_type", "swiglu")
```

```
191
192     if ffn_type == "silu":
193         self.ffn = SiLU(d_model, d_ff, device, dtype)
194     elif ffn_type == "swiglu":
195         self.ffn = SwiGLU(d_model, d_ff, device, dtype)
196     else:
197         raise ValueError(f"Unsupported ffn_type: {ffn_type}")
198
199     def forward(self, x: torch.Tensor):
200         x = x + self.attn(self.ln1(x), self.rope)
201         x = x + self.ffn(self.ln2(x))
202         return x
203
204
205 class Embedding(torch.nn.Module):
206     def __init__(
207         self,
208         num_embeddings: int,
209         embedding_dim: int,
210         device: torch.device | None = None,
211         dtype: torch.dtype | None = None,
212         **kwargs,
213     ):
214         super().__init__()
215
216         mean = 0
217         std = 1
218         lower = -3
219         upper = 3
220
221         if kwargs.get("embedding_std", None) is not None:
222             std = kwargs.get("embedding_std")
223
224         w = torch.empty((num_embeddings, embedding_dim), device=device, dtype=dtype)
225         torch.nn.init.trunc_normal_(w, mean=mean, std=std, a=lower, b=upper)
226
227         self.weight = torch.nn.Parameter(w)
228
229     def forward(self, token_ids: torch.Tensor) -> torch.Tensor:
230         return self.weight[token_ids]
231
232
233 class Transformer(torch.nn.Module):
234     def __init__(
235         self,
236         d_model: int,
237         num_heads: int,
238         d_ff: int,
239         vocab_size: int,
240         context_length: int,
```

```

241         num_layers: int,
242         rope_theta: float = 10000.0,
243         device=None,
244         dtype=None,
245         **kwargs,
246     ):
247         super().__init__()
248
249         self.context_length = context_length
250         self.token_embeddings = Embedding(vocab_size, d_model, device, dtype,
**kwargs)
251
252         if d_model % num_heads != 0:
253             raise ValueError("d_model must be divisible by num_heads")
254
255         d_head = d_model // num_heads
256         rope = RotaryPositionalEmbedding(rope_theta, d_head, context_length,
device=device, dtype=dtype)
257
258         self.layers = torch.nn.ModuleList(
259             [Block(d_model, num_heads, d_ff, rope, device, dtype, **kwargs) for _ in
range(num_layers)]
260         )
261
262         self.ln_final = RMSNorm(d_model, device=device, dtype=dtype)
263         self.lm_head = Linear(d_model, vocab_size, device, dtype)
264
265         if kwargs.get("weight_tying", False):
266             self.lm_head.weight = self.token_embeddings.weight
267
268     def forward(self, x: torch.Tensor) -> torch.Tensor:
269         batch_size, seq_len = x.shape
270
271         if seq_len > self.context_length:
272             raise ValueError(f"Input sequence length ({seq_len}) exceeds model
context length ({self.context_length})")
273
274         x = self.token_embeddings(x)
275
276         for layer in self.layers:
277             x = layer(x)
278
279         x = self.ln_final(x)
280         x = self.lm_head(x)
281
282         return x
283

```

## 4 Training a Transformer LM

We now have the steps to preprocess the data (via tokenizer) and the model (Transformer). What remains is to build all of the code to support training. This consists of the following:

- **Loss:** we need to define the loss function (cross-entropy).
- **Optimizer:** we need to define the optimizer to minimize this loss (AdamW).
- **Training loop:** we need all the supporting infrastructure that loads data, saves checkpoints, and manages training.

### 4.1 Cross-entropy loss

Recall that the Transformer language model defines a distribution  $p_\theta(x_{i+1} \mid x_{1:i})$  for each sequence  $x$  of length  $m+1$  and  $i = 1, \dots, m$ . Given a training set  $D$  consisting of sequences of length  $m$ , we define the standard cross-entropy (negative log-likelihood) loss function:

$$\ell(\theta; D) = \frac{1}{|D|m} \sum_{x \in D} \sum_{i=1}^m -\log p_\theta(x_{i+1} \mid x_{1:i}). \quad (16)$$

(Note that a single forward pass in the Transformer yields  $p_\theta(x_{i+1} \mid x_{1:i})$  for all  $i = 1, \dots, m$ .)

In particular, the Transformer computes logits  $o_i \in \mathbb{R}^{\text{vocab\_size}}$  for each position  $i$ , which results in:<sup>6</sup>

$$p(x_{i+1} \mid x_{1:i}) = \text{softmax}(o_i)[x_{i+1}] = \frac{\exp(o_i[x_{i+1}])}{\sum_{a=1}^{\text{vocab\_size}} \exp(o_i[a])}. \quad (17)$$

The cross entropy loss is generally defined with respect to the vector of logits  $o_i \in \mathbb{R}^{\text{vocab\_size}}$  and target  $x_{i+1}$ .<sup>7</sup>

Implementing the cross entropy loss requires some care with numerical issues, just like in the case of softmax.

#### Problem (cross\_entropy): Implement Cross entropy

**Deliverable:** Write a function to compute the cross entropy loss, which takes in predicted logits ( $o_i$ ) and targets ( $x_{i+1}$ ) and computes the cross entropy  $\ell_i = -\log \text{softmax}(o_i)[x_{i+1}]$ . Your function should handle the following:

- Subtract the largest element for numerical stability.
- Cancel out log and exp whenever possible.
- Handle any additional batch dimensions and return the *average across the batch*. As with section 3.3, we assume batch-like dimensions always come first, before the vocabulary size dimension.

Implement `[adapters.run_cross_entropy]`, then run `uv run pytest -k test_cross_entropy` to test your implementation.

**Perplexity** Cross entropy suffices for training, but when we evaluate the model, we also want to report perplexity. For a sequence of length  $m$  where we suffer cross-entropy losses  $\ell_1, \dots, \ell_m$ :

$$\text{perplexity} = \exp \left( \frac{1}{m} \sum_{i=1}^m \ell_i \right). \quad (18)$$

<sup>6</sup>Note that  $o_i[k]$  refers to value at index  $k$  of the vector  $o_i$ .

<sup>7</sup>This corresponds to the cross entropy between the Dirac delta distribution over  $x_{i+1}$  and the predicted  $\text{softmax}(o_i)$  distribution.

## 4.2 The SGD Optimizer

Now that we have a loss function, we will begin our exploration of optimizers. The simplest gradient-based optimizer is Stochastic Gradient Descent (SGD). We start with randomly initialized parameters  $\theta_0$ . Then for each step  $t = 0, \dots, T - 1$ , we perform the following update:

$$\theta_{t+1} \leftarrow \theta_t - \alpha_t \nabla L(\theta_t; B_t), \quad (19)$$

where  $B_t$  is a random batch of data sampled from the dataset  $D$ , and the *learning rate*  $\alpha_t$  and *batch size*  $|B_t|$  are hyperparameters.

### 4.2.1 Implementing SGD in PyTorch

To implement our optimizers, we will subclass the PyTorch `torch.optim.Optimizer` class. An `Optimizer` subclass must implement two methods:

`def __init__(self, params, ...)` should initialize your optimizer. Here, `params` will be a collection of parameters to be optimized (or parameter groups, in case the user wants to use different hyperparameters, such as learning rates, for different parts of the model). **Make sure to pass `params` to the `__init__` method of the base class, which will store these parameters for use in `step`.** You can take additional arguments depending on the optimizer (e.g., the learning rate is a common one), and pass them to the base class constructor as a dictionary, where keys are the names (strings) you choose for these parameters.

`def step(self)` should make one update of the parameters. During the training loop, this will be called after the backward pass, so you have access to the gradients on the last batch. This method should iterate through each parameter tensor `p` and modify them *in place*, i.e. setting `p.data`, which holds the tensor associated with that parameter based on the gradient `p.grad` (if it exists), the tensor representing the gradient of the loss with respect to that parameter.

The PyTorch optimizer API has a few subtleties, so it's easier to explain it with an example. To make our example richer, we'll implement a slight variation of SGD where the learning rate decays over training, starting with an initial learning rate  $\alpha$  and taking successively smaller steps over time:

$$\theta_{t+1} = \theta_t - \frac{\alpha}{\sqrt{t+1}} \nabla L(\theta_t; B_t) \quad (20)$$

Let's see how this version of SGD would be implemented as a PyTorch `Optimizer`:

```
from collections.abc import Callable, Iterable
from typing import Optional
import torch
import math

class SGD(torch.optim.Optimizer):
    def __init__(self, params, lr=1e-3):
        if lr < 0:
            raise ValueError(f"Invalid learning rate: {lr}")
        defaults = {"lr": lr}
        super().__init__(params, defaults)

    def step(self, closure: Optional[Callable] = None):
        loss = None if closure is None else closure()
        for group in self.param_groups:
            lr = group["lr"] # Get the learning rate.
```

this is inherited  
from torch.optim.Optimizer

```

for p in group["params"]:
    if p.grad is None:
        continue

    state = self.state[p] # Get state associated with p.
    t = state.get("t", 0) # Get iteration number from the state, or initial value.
    grad = p.grad.data # Get the gradient of loss with respect to p.
    p.data -= lr / math.sqrt(t + 1) * grad # Update weight tensor in-place.
    state["t"] = t + 1 # Increment iteration number.

return loss

```

In `__init__`, we pass the parameters to the optimizer, as well as default hyperparameters, to the base class constructor (the parameters might come in groups, each with different hyperparameters). In case the parameters are just a single collection of `torch.nn.Parameter` objects, the base constructor will create a single group and assign it the default hyperparameters. Then, in `step`, we iterate over each parameter group, then over each parameter in that group, and apply Eq 20. Here, we keep the iteration number as a state associated with each parameter: we first read this value, use it in the gradient update, and then update it. The API specifies that the user might pass in a callable `closure` to re-compute the loss before the optimizer step. We won't need this for the optimizers we'll use, but we add it to comply with the API.

To see this working, we can use the following minimal example of a *training loop*:

```

weights = torch.nn.Parameter(5 * torch.randn((10, 10)))
opt = SGD([weights], lr=1)

for t in range(100):
    opt.zero_grad() # Reset the gradients for all learnable parameters.
    loss = (weights**2).mean() # Compute a scalar loss value.
    print(loss.cpu().item())
    loss.backward() # Run backward pass, which computes gradients.
    opt.step() # Run optimizer step.

```

This is the typical structure of a training loop: in each iteration, we will compute the loss and run a step of the optimizer. When training language models, our learnable parameters will come from the model (in PyTorch, `m.parameters()` gives us this collection). The loss will be computed over a sampled batch of data, but the basic structure of the training loop will be the same.

#### Problem (learning\_rate\_tuning): Tuning the learning rate (1 point)

As we will see, one of the hyperparameters that affects training the most is the learning rate. Let's see that in practice in our toy example. Run the SGD example above with three other values for the learning rate: `1e1`, `1e2`, and `1e3`, for just 10 training iterations. What happens with the loss for each of these learning rates? Does it decay faster, slower, or does it diverge (i.e., increase over the course of training)?

**Deliverable:** A one-two sentence response with the behaviors you observed.

### 4.3 AdamW

Modern language models are typically trained with more sophisticated optimizers, instead of SGD. Most optimizers used recently are derivatives of the Adam optimizer [Kingma and Ba, 2015]. We will use AdamW [Loshchilov and Hutter, 2019], which is in wide use in recent work. AdamW proposes a modification to Adam that improves regularization by adding *weight decay* (at each iteration, we pull the parameters towards 0),



in a way that is decoupled from the gradient update. We will implement AdamW as described in algorithm 2 of [Loshchilov and Hutter, 2019].

AdamW is *stateful*: for each parameter, it keeps track of a running estimate of its first and second moments. Thus, AdamW uses additional memory in exchange for improved stability and convergence. Besides the learning rate  $\alpha$ , AdamW has a pair of hyperparameters  $(\beta_1, \beta_2)$  that control the updates to the moment estimates, and a weight decay rate  $\lambda$ . Typical applications set  $(\beta_1, \beta_2)$  to  $(0.9, 0.999)$ , but large language models like LLaMA [Touvron et al., 2023] and GPT-3 [Brown et al., 2020] are often trained with  $(0.9, 0.95)$ . The algorithm can be written as follows, where  $\epsilon$  is a small value (e.g.,  $10^{-8}$ ) used to improve numerical stability in case we get extremely small values in  $v$ :

---

**Algorithm 1** AdamW Optimizer

---

```

init( $\theta$ ) (Initialize learnable parameters)
 $m \leftarrow 0$  (Initial value of the first moment vector; same shape as  $\theta$ )
 $v \leftarrow 0$  (Initial value of the second moment vector; same shape as  $\theta$ )
for  $t = 1, \dots, T$  do
    Sample batch of data  $B_t$ 
     $g \leftarrow \nabla_{\theta} \ell(\theta; B_t)$  (Compute the gradient of the loss at the current time step)
     $m \leftarrow \beta_1 m + (1 - \beta_1)g$  (Update the first moment estimate)
     $v \leftarrow \beta_2 v + (1 - \beta_2)g^2$  (Update the second moment estimate)
     $\alpha_t \leftarrow \alpha \frac{\sqrt{1 - (\beta_2)^t}}{1 - (\beta_1)^t}$  (Compute adjusted  $\alpha$  for iteration  $t$ )
     $\theta \leftarrow \theta - \alpha_t \frac{m}{\sqrt{v} + \epsilon}$  (Update the parameters)
     $\theta \leftarrow \theta - \alpha \lambda \theta$  (Apply weight decay)
end for

```

---

Note that  $t$  starts at 1. You will now implement this optimizer.

**Problem (adamw): Implement AdamW (2 points)**

---

**Deliverable:** Implement the AdamW optimizer as a subclass of `torch.optim.Optimizer`. Your class should take the learning rate  $\alpha$  in `__init__`, as well as the  $\beta$ ,  $\epsilon$  and  $\lambda$  hyperparameters. To help you keep state, the base `Optimizer` class gives you a dictionary `self.state`, which maps `nn.Parameter` objects to a dictionary that stores any information you need for that parameter (for AdamW, this would be the moment estimates). Implement `[adapters.get_adamw_cls]` and make sure it passes `uv run pytest -k test_adamw`.

**Problem (adamwAccounting): Resource accounting for training with AdamW (2 points)**

---

Let us compute how much memory and compute running AdamW requires. Assume we are using float32 for every tensor.

- (a) How much peak memory does running AdamW require? Decompose your answer based on the memory usage of the parameters, activations, gradients, and optimizer state. Express your answer in terms of the `batch_size` and the model hyperparameters (`vocab_size`, `context_length`, `num_layers`, `d_model`, `num_heads`). Assume `d_ff = 4 × d_model`.

For simplicity, when calculating memory usage of activations, consider only the following components:

- Transformer block
  - RMSNorm(s)

- Multi-head self-attention sublayer:  $QKV$  projections,  $Q^\top K$  matrix multiply, softmax, weighted sum of values, output projection.
- Position-wise feed-forward:  $W_1$  matrix multiply, SiLU,  $W_2$  matrix multiply
- final RMSNorm
- output embedding
- cross-entropy on logits

**Deliverable:** An algebraic expression for each of parameters, activations, gradients, and optimizer state, as well as the total.

- (b) Instantiate your answer for a GPT-2 XL-shaped model to get an expression that only depends on the `batch_size`. What is the maximum batch size you can use and still fit within 80GB memory?

**Deliverable:** An expression that looks like  $a \cdot \text{batch\_size} + b$  for numerical values  $a, b$ , and a number representing the maximum batch size.

- (c) How many FLOPs does running one step of AdamW take?

**Deliverable:** An algebraic expression, with a brief justification.

- (d) Model FLOPs utilization (MFU) is defined as the ratio of observed throughput (tokens per second) relative to the hardware's theoretical peak FLOP throughput [Chowdhery et al., 2022]. An NVIDIA A100 GPU has a theoretical peak of 19.5 teraFLOP/s for float32 operations. Assuming you are able to get 50% MFU, how long would it take to train a GPT-2 XL for 400K steps and a batch size of 1024 on a single A100? Following Kaplan et al. [2020] and Hoffmann et al. [2022], assume that the backward pass has twice the FLOPs of the forward pass.

**Deliverable:** The number of days training would take, with a brief justification.

## 4.4 Learning rate scheduling

The value for the learning rate that leads to the quickest decrease in loss often varies during training. In training Transformers, it is typical to use a learning rate *schedule*, where we start with a bigger learning rate, making quicker updates in the beginning, and slowly decay it to a smaller value as the model trains. In this assignment, we will implement the *cosine annealing schedule* used to train LLaMA [Touvron et al., 2023].

A scheduler is simply a function that takes the current step  $t$  and other relevant parameters (such as the initial and final learning rates), and returns the learning rate to use for the gradient update at step  $t$ . The simplest schedule is the constant function, which will return the same learning rate given any  $t$ .

The cosine annealing learning rate schedule takes (i) the current iteration  $t$ , (ii) the maximum learning rate  $\alpha_{\max}$ , (iii) the minimum (final) learning rate  $\alpha_{\min}$ , (iv) the number of *warm-up* iterations  $T_w$ , and (v) the number of cosine annealing iterations  $T_c$ . The learning rate at iteration  $t$  is defined as:

**(Warm-up)** If  $t < T_w$ , then  $\alpha_t = \frac{t}{T_w} \alpha_{\max}$ .

**(Cosine annealing)** If  $T_w \leq t \leq T_c$ , then  $\alpha_t = \alpha_{\min} + \frac{1}{2} \left( 1 + \cos \left( \frac{t - T_w}{T_c - T_w} \pi \right) \right) (\alpha_{\max} - \alpha_{\min})$ .

**(Post-annealing)** If  $t > T_c$ , then  $\alpha_t = \alpha_{\min}$ .

<sup>8</sup>It's sometimes common to use a schedule where the learning rate rises back up (restarts) to help get past local minima.

**Problem (learning\_rate\_schedule): Implement cosine learning rate schedule with warmup**

---

Write a function that takes  $t$ ,  $\alpha_{\max}$ ,  $\alpha_{\min}$ ,  $T_w$  and  $T_c$ , and returns the learning rate  $\alpha_t$  according to the scheduler defined above. Then implement `[adapters.get_lr_cosine_schedule]` and make sure it passes `uv run pytest -k test_get_lr_cosine_schedule`.

## 4.5 Gradient clipping

During training, we can sometimes hit training examples that yield large gradients, which can destabilize training. To mitigate this, one technique often employed in practice is *gradient clipping*. The idea is to enforce a limit on the norm of the gradient after each backward pass before taking an optimizer step.

Given the gradient (for all parameters)  $g$ , we compute its  $\ell_2$ -norm  $\|g\|_2$ . If this norm is less than a maximum value  $M$ , then we leave  $g$  as is; otherwise, we scale  $g$  down by a factor of  $\frac{M}{\|g\|_2 + \epsilon}$  (where a small  $\epsilon$ , like  $10^{-6}$ , is added for numeric stability). Note that the resulting norm will be just under  $M$ .

**Problem (gradient\_clipping): Implement gradient clipping (1 point)**

---

Write a function that implements gradient clipping. Your function should take a list of parameters and a maximum  $\ell_2$ -norm. It should modify each parameter gradient in place. Use  $\epsilon = 10^{-6}$  (the PyTorch default). Then, implement the adapter `[adapters.run_gradient_clipping]` and make sure it passes `uv run pytest -k test_gradient_clipping`.

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snider/cs336\_basics/loss.py

```
1 import torch
2
3
4 def cross_entropy_loss_naive(logits: torch.Tensor, targets: torch.Tensor) ->
  torch.Tensor:
5     """
6     Computes the cross entropy loss between logits and target indices.
7
8     Args:
9         logits (torch.Tensor): Logits with shape (... , vocab_size)
10        targets (torch.Tensor): Target indices with shape (...)
11
12    Returns:
13        torch.Tensor: Average cross entropy loss across the batch
14    """
15    max_logits = logits.max(dim=-1, keepdim=True).values
16    logits_shifted = logits - max_logits
17    sum_exp = torch.exp(logits_shifted).sum(dim=-1, keepdim=True)
18    log_sum_exp = torch.log(sum_exp)
19    log_probs = logits_shifted - log_sum_exp
20
21    target_log_probs = log_probs.gather(dim=-1,
index=targets.unsqueeze(-1)).squeeze(-1)
22
23    return -target_log_probs.mean()
24
25
26 @torch.compile
27 def cross_entropy_loss(logits: torch.Tensor, targets: torch.Tensor) -> torch.Tensor:
28     max_vals, _ = logits.max(dim=-1, keepdim=True)
29     target_logits = logits.gather(dim=-1, index=targets.unsqueeze(-1))
30     shifted = logits - max_vals
31     sum_exp = shifted.exp().sum(dim=-1, keepdim=True)
32     log_sum_exp = sum_exp.log()
33     return -((target_logits - max_vals) - log_sum_exp).mean()
34
```

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snider/cs336\_basics/adamw.py

```

1 import torch
2 import math
3
4
5 class AdamW(torch.optim.Optimizer):
6     def __init__(
7         self,
8         params,
9         lr: float = 1e-3,
10        betas: tuple[float, float] = (0.9, 0.95),
11        eps: float = 1e-8,
12        weight_decay: float = 0.1,
13        **kwargs,
14    ):
15        defaults = {"lr": lr, "betas": betas, "eps": eps, "weight_decay":
weight_decay}
16        super().__init__(params, defaults)
17
18        self.lr = lr
19        self.betas = betas
20        self.eps = eps
21        self.weight_decay = weight_decay
22
23    @torch.no_grad()
24    def step(self, closure=None):
25        loss = None if closure is None else closure()
26
27        for group in self.param_groups:
28            b1, b2 = group["betas"]
29            lr = group["lr"]
30            eps = group["eps"]
31            wd = group["weight_decay"]
32
33            for p in group["params"]:
34                if p.grad is None:
35                    continue
36
37                state = self.state[p]
38
39                m = state.get("m", torch.zeros_like(p.data))
40                v = state.get("v", torch.zeros_like(p.data))
41                t = state.get("t", 1)
42
43                grad = p.grad
44
45                state["m"] = b1 * m + (1 - b1) * grad
46                state["v"] = b2 * v + (1 - b2) * grad.pow(2)
47

```

```
48         step_size = lr * (math.sqrt(1 - b2**t) / (1 - b1**t))
49
50         # must use addcddiv to update in place! otherwise extra memory is
allotted
51         p.data.addcddiv_(state["m"], torch.sqrt(state["v"]) + eps, value=-
step_size)
52
53         if wd != 0:
54             p.data.add_(p.data, alpha=-lr * wd)
55
56         state["t"] = t + 1
57
58     return loss
59
```

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snyder/cs336\_basics/lr\_schedule.py

```

1  import math
2
3  from torch.optim import Optimizer
4
5
6  class LRSchedule:
7      def __init__(
8          self,
9          lr_max: float,
10         warmup_iters: int,
11         schedule: list[dict],
12         optimizer: Optimizer | None = None,
13         param_groups: list[dict] | None = None,
14     ):
15         """
16         A steppable learning rate scheduler supporting arbitrary multi-phase decay
17         schedules.
18
19         Args:
20             lr_max: The maximum learning rate.
21             warmup_iters: The number of iterations for linear warmup from zero to
22             lr_max
23             schedule: A list of dictionaries, each containing the following keys:
24                 - until_iter: The iteration at which the phase should end
25                 - to_lr: The learning rate at which the phase should end
26                 - type: The type of decay to use for the phase (linear, cosine, exp)
27             optimizer (optional): The optimizer whose param groups will be updated
28             param_groups (optional): The param groups to update the learning rate for
29                 - If provided, will use these param groups instead of
30             optimizer.param_groups
31         """
32         self.lr_max = lr_max
33         self.warmup_iters = warmup_iters
34         self.schedule = schedule
35         self.optimizer = optimizer
36         self.param_groups = param_groups
37         self.reset()
38
39     def step(self):
40         self.it += 1
41         self.lr = lr_schedule(self.it, self.lr_max, self.warmup_iters, self.schedule)
42
43         param_groups = self.param_groups if self.param_groups is not None else
44         self.optimizer.param_groups
45
46         for param_group in param_groups:
47             param_group["lr"] = self.lr

```

```

45     def reset(self):
46         self.it = -1
47         self.step()
48
49     def __repr__(self):
50         return f"LRSchedule(lr_max={self.lr_max}, warmup_iters={self.warmup_iters},
schedule={self.schedule}, optimizer={self.optimizer}, param_groups=
{self.param_groups})"
51
52
53 def lr_schedule(
54     it: int,
55     lr_max: float,
56     warmup_iters: int,
57     schedule: list[dict],
58 ):
59     """A learning rate scheduler supporting arbitrary multi-phase decay schedules.
60
61     Args:
62         it: The current iteration.
63         lr_max: The maximum learning rate.
64         warmup_iters: The number of iterations for linear warmup from zero to lr_max
65         schedule: A list of dictionaries, each containing the following keys:
66             - until_iter: The iteration at which the phase should end
67             - to_lr: The learning rate at which the phase should end
68             - type: The type of decay to use for the phase (linear, cosine, exp)
69
70     """
71     if it < warmup_iters:
72         return (it / warmup_iters) * lr_max
73
74     phase_max_lr = lr_max
75     phase_start_iter = warmup_iters
76
77     for phase in schedule:
78         phase_min_lr = phase["to_lr"]
79
80         if it <= phase["until_iter"]:
81             decay_step = it - phase_start_iter
82             decay_steps = phase["until_iter"] - phase_start_iter
83
84             if phase["type"] == "linear":
85                 return phase_max_lr - (decay_step / decay_steps) * (phase_max_lr -
phase_min_lr)
86             elif phase["type"] == "cosine":
87                 cos = math.cos((decay_step / decay_steps) * math.pi)
88                 return phase_min_lr + 1 / 2 * (1 + cos) * (phase_max_lr -
phase_min_lr)
89             elif phase["type"] == "exp":
90                 return phase_max_lr * (phase_min_lr / phase_max_lr) ** (decay_step /
decay_steps)

```



```
91     phase_start_iter = phase["until_iter"]
92     phase_max_lr = phase["to_lr"]
93
94     return schedule[-1]["to_lr"]
95
96
97 def lr_linear_schedule(it: int, lr_max: float, lr_min: float, warmup_iters: int,
98 linear_cycle_iters: int):
99     if it < warmup_iters:
100         return (it / warmup_iters) * lr_max
101
102     if it <= linear_cycle_iters:
103         decay_step = it - warmup_iters
104         decay_steps = linear_cycle_iters - warmup_iters
105         return lr_max - (decay_step / decay_steps) * (lr_max - lr_min)
106
107     return lr_min
108
109 def lr_cosine_schedule(it: int, lr_max: float, lr_min: float, warmup_iters: int,
110 cosine_cycle_iters: int):
111     if it < warmup_iters:
112         return (it / warmup_iters) * lr_max
113
114     if it <= cosine_cycle_iters:
115         decay_step = it - warmup_iters
116         decay_steps = cosine_cycle_iters - warmup_iters
117         cos = math.cos((decay_step / decay_steps) * math.pi)
118         return lr_min + 1 / 2 * (1 + cos) * (lr_max - lr_min)
119
120     return lr_min
121
122 def lr_double_schedule(
123     it: int,
124     lr_max: float,
125     lr_inter: int,
126     lr_min: float,
127     warmup_iters: int,
128     phase_one_iters: int,
129     phase_two_iters: int,
130     phase_two_type: str,
131 ):
132     """
133     A double-decay learning rate schedule.
134
135     Args:
136         it: The current iteration.
137         lr_max: Max. LR, to which we warm up linearly.
138         lr_inter: LR to which we decay exponentially from lr_max.
139         lr_min: Min. LR, to which we decay from lr_inter, linearly or cosine.
```

```

140     warmup_iters: The number of iters for linear warmup from zero to lr_max
141     exp_decay_iters: The iter at which the exponential decay phase should end
142     phase_two_iters: The iter at which the second decay phase (linear or cosine)
should end
143     phase_two_type: The type of decay to use for the second phase (linear or
cosine)
144
145     Note:
146     - exp_decay_iters is NOT the number of iterations for the exponential decay
phase.
147         It is the iter at which the exponential decay should end.
148     - phase_two_iters is NOT the number of iterations for the second decay phase.
149         It is the iter at which the second decay should end.
150     Example:
151     - Want: warmup for 1000 iters, exp decay for 1000 iters, linear decay for
1000 iters
152     - Set:
153         warmup_iters = 1000
154         exp_decay_iters = 2000
155         phase_two_iters = 3000
156         phase_two_type = "linear"
157     """
158     if it < warmup_iters:
159         # We're in the warmup phase
160         return (it / warmup_iters) * lr_max
161
162     if it <= phase_one_iters:
163         # We're in the exponential decay phase
164         decay_step = it - warmup_iters
165         decay_steps = phase_one_iters - warmup_iters
166         return lr_max * (lr_inter / lr_max) ** (decay_step / decay_steps)
167
168     # We're in phase two (linear or cosine decay from lr_inter to lr_min)
169     it2 = it - phase_one_iters
170     phase_two_decay_steps = phase_two_iters - phase_one_iters
171
172     if phase_two_type == "linear":
173         # The second decay phase of the schedule is linear
174         return lr_linear_schedule(
175             it2, lr_max=lr_inter, lr_min=lr_min, warmup_iters=0,
linear_cycle_iters=phase_two_decay_steps
176         )
177
178     if phase_two_type == "cosine":
179         # The second decay phase of the schedule is cosine
180         return lr_cosine_schedule(
181             it2, lr_max=lr_inter, lr_min=lr_min, warmup_iters=0,
cosine_cycle_iters=phase_two_decay_steps
182         )
183
184     return lr_min

```

```
185
186
187 def seq_len_schedule(
188     it: int,
189     seq_len_min: int,
190     schedule: list[dict],
191 ):
192     """A sequence length scheduler supporting arbitrary multi-phase decay schedules.
193
194     Args:
195         it: The current iteration.
196         seq_len_min: The minimum sequence length.
197         schedule: A list of dictionaries, each containing the following keys:
198             - until_iter: The iteration at which the phase should end
199             - to_seq_len: The sequence length at which the phase should end
200     """
201     phase_min_seq_len = seq_len_min
202     phase_start_iter = 0
203
204     for phase in schedule:
205         phase_max_seq_len = phase["to_seq_len"]
206
207         if it <= phase["until_iter"]:
208             growth_step = it - phase_start_iter
209             growth_steps = phase["until_iter"] - phase_start_iter
210
211             return int(phase_min_seq_len + (growth_step / growth_steps) *
212 (phase_max_seq_len - phase_min_seq_len))
213
214         phase_start_iter = phase["until_iter"]
215         phase_min_seq_len = phase["to_seq_len"]
216
217     return schedule[-1]["to_seq_len"]
218
219 def batch_size_schedule(it: int, batch_size_max: int, schedule: list[dict]):
220     """A batch size scheduler supporting arbitrary multi-phase decay schedules.
221
222     Args:
223         it: The current iteration.
224         batch_size_max: The maximum batch size.
225         schedule: A list of dictionaries, each containing the following keys:
226             - until_iter: The iteration at which the phase should end
227             - to_batch_size: The batch size at which the phase should end
228     """
229     phase_max_batch_size = batch_size_max
230     phase_start_iter = 0
231
232     for phase in schedule:
233         phase_min_batch_size = phase["to_batch_size"]
```

```
234
235     if it <= phase["until_iter"]:
236         decay_step = it - phase_start_iter
237         decay_steps = phase["until_iter"] - phase_start_iter
238
239         return int(
240             phase_max_batch_size - (decay_step / decay_steps) *
241             (phase_max_batch_size - phase_min_batch_size)
242         )
243
244     phase_start_iter = phase["until_iter"]
245     phase_max_batch_size = phase["to_batch_size"]
246
247     return schedule[-1]["to_batch_size"]
```

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snyder/cs336\_basics/gradient\_clip.py

```
1 from collections.abc import Iterable
2 import torch
3
4
5 @torch.compile
6 def gradient_clip(parameters: Iterable[torch.nn.Parameter], max_l2_norm: float):
7     grads = [p.grad for p in parameters if p.grad is not None]
8
9     if not grads:
10         return torch.tensor(0.0)
11
12     stacked_grads = torch.stack([torch.norm(g.detach(), p=2) for g in grads])
13     total_norm = torch.norm(stacked_grads, p=2)
14
15     if total_norm > max_l2_norm:
16         scale = max_l2_norm / (total_norm + 1e-6)
17         for grad in grads:
18             grad.detach().mul_(scale)
19
20     return total_norm
21
```

## 5 Training loop

We will now finally put together the major components we’ve built so far: the tokenized data, the model, and the optimizer.

### 5.1 Data Loader

The tokenized data (e.g., that you prepared in `tokenizer_experiments`) is a single sequence of tokens  $x = (x_1, \dots, x_n)$ . Even though the source data might consist of separate documents (e.g., different web pages, or source code files), a common practice is to concatenate all of those into a single sequence of tokens, adding a delimiter between them (such as the `<|endoftext|>` token).

A *data loader* turns this into a stream of *batches*, where each batch consists of  $B$  sequences of length  $m$ , paired with the corresponding next tokens, also with length  $m$ . For example, for  $B = 1, m = 3$ ,  $([x_2, x_3, x_4], [x_3, x_4, x_5])$  would be one potential batch.

Loading data in this way simplifies training for a number of reasons. First, any  $1 \leq i < n - m$  gives a valid training sequence, so sampling sequences are trivial. **Since all training sequences have the same length, there’s no need to pad input sequences, which improves hardware utilization** (also by increasing batch size  $B$ ). Finally, we also don’t need to fully load the full dataset to sample training data, making it easy to handle large datasets that might not otherwise fit in memory.

#### Problem (data\_loading): Implement data loading (2 points)

**Deliverable:** Write a function that takes a numpy array  $x$  (integer array with token IDs), a `batch_size`, a `context_length` and a PyTorch device string (e.g., `'cpu'` or `'cuda:0'`), and returns a pair of tensors: the sampled input sequences and the corresponding next-token targets. Both tensors should have shape  $(\text{batch\_size}, \text{context\_length})$  containing token IDs, and both should be placed on the requested device. To test your implementation against our provided tests, you will first need to implement the test adapter at `[adapters.run_get_batch]`. Then, run `uv run pytest -k test_get_batch` to test your implementation.

#### Low-Resource/Downscaling Tip: Data loading on CPU or Apple Silicon

If you are planning to train your LM on CPU or Apple Silicon, you need to move your data to the correct device (and similarly, you should use the same device for your model later on).

If you are on CPU, you can use the `'cpu'` device string, and on Apple Silicon (M\* chips), you can use the `'mps'` device string.

For more on MPS, checkout these resources:

- <https://developer.apple.com/metal/pytorch/>
- <https://pytorch.org/docs/main/notes/mps.html>

What if the dataset is too big to load into memory? We can use a Unix systemcall named `mmap` which maps a file on disk to virtual memory, and lazily loads the file contents when that memory location is accessed. Thus, you can “pretend” you have the entire dataset in memory. Numpy implements this through `np.memmap` (or the flag `mmap_mode='r'` to `np.load`, if you originally saved the array with `np.save`), which will return a numpy array-like object that loads the entries on-demand as you access them. **When sampling from your dataset (i.e., a numpy array) during training, be sure load the dataset in memory-mapped mode** (via `np.memmap` or the flag `mmap_mode='r'` to `np.load`, depending on how you saved the array). Make sure you also specify a `dtype` that matches the array that you’re loading. It may be helpful to explicitly verify that the memory-mapped data looks correct (e.g., doesn’t contain values beyond the expected vocabulary size).

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snyder/cs336\_basics/data\_loader.py

```

1 import numpy as np
2 import torch
3
4
5 def get_batch(x: np.ndarray, batch_size: int, context_length: int, device: str):
6     """
7     Takes a numpy array of token IDs, batches them, and returns a pair of tensors:
8     - (batch_size, context_length) – the actual batch
9     - (batch_size, context_length) – the next token ID for each sample in the batch
10
11     Args:
12         x: np.ndarray – integer array of training data
13         batch_size: int – number of samples per batch
14         context_length: int – length of each sequence in batch
15         device: str – device to load the data on
16     """
17     # Maximum valid starting index to ensure we have enough tokens for a full sequence
    plus one
18     max_start_idx = len(x) - context_length - 1
19
20     if max_start_idx < 0:
21         raise ValueError(f"Input array length {len(x)} is too short for context_length {context_length}")
22
23     # Sample batch_size random starting positions all at once
24     start_indices = np.random.randint(0, max_start_idx + 1, size=batch_size)
25
26     # Prepare arrays to hold our sequences
27     x_sequences = np.zeros((batch_size, context_length), dtype=np.int64)
28     y_sequences = np.zeros((batch_size, context_length), dtype=np.int64)
29
30     # Fill the arrays with the appropriate sequences
31     for i, start_idx in enumerate(start_indices):
32         x_sequences[i] = x[start_idx : start_idx + context_length]
33         y_sequences[i] = x[start_idx + 1 : start_idx + context_length + 1]
34
35     x_batch = torch.from_numpy(x_sequences)
36     y_batch = torch.from_numpy(y_sequences)
37
38     if device.startswith("cuda"):
39         # Pin memory if on GPU
40         x_batch, y_batch = (
41             x_batch.pin_memory().to(device, non_blocking=True),
42             y_batch.pin_memory().to(device, non_blocking=True),
43         )
44     else:
45         x_batch, y_batch = x_batch.to(device), y_batch.to(device)

```

```
46  
47     return x_batch, y_batch  
48
```



## 5.2 Checkpointing

In addition to loading data, we will also need to save models as we train. When running jobs, we often want to be able to resume a training run that for some reason stopped midway (e.g., due to your job timing out, machine failure, etc). Even when all goes well, we might also want to later have access to intermediate models (e.g., to study training dynamics post-hoc, take samples from models at different stages of training, etc).

A checkpoint should have all the states that we need to resume training. We of course want to be able to restore model weights at a minimum. If using a stateful optimizer (such as AdamW), we will also need to save the optimizer's state (e.g., in the case of AdamW, the moment estimates). Finally, to resume the learning rate schedule, we will need to know the iteration number we stopped at. PyTorch makes it easy to save all of these: every `nn.Module` has a `state_dict()` method that returns a dictionary with all learnable weights; we can restore these weights later with the sister method `load_state_dict()`. The same goes for any `nn.optim.Optimizer`. Finally, `torch.save(obj, dest)` can dump an object (e.g., a dictionary containing tensors in some values, but also regular Python objects like integers) to a file (path) or file-like object, which can then be loaded back into memory with `torch.load(src)`.

### Problem (checkpointing): Implement model checkpointing (1 point)

Implement the following two functions to load and save checkpoints:

`def save_checkpoint(model, optimizer, iteration, out)` should dump all the state from the first three parameters into the file-like object `out`. You can use the `state_dict` method of both the model and the optimizer to get their relevant states and use `torch.save(obj, out)` to dump `obj` into `out` (PyTorch supports either a path or a file-like object here). A typical choice is to have `obj` be a dictionary, but you can use whatever format you want as long as you can load your checkpoint later.

This function expects the following parameters:

```
model: torch.nn.Module
optimizer: torch.optim.Optimizer
iteration: int
out: str | os.PathLike | typing.BinaryIO | typing.IO[bytes]
```

`def load_checkpoint(src, model, optimizer)` should load a checkpoint from `src` (path or file-like object), and then recover the model and optimizer states from that checkpoint. Your function should return the iteration number that was saved to the checkpoint. You can use `torch.load(src)` to recover what you saved in your `save_checkpoint` implementation, and the `load_state_dict` method in both the model and optimizers to return them to their previous states.

This function expects the following parameters:

```
src: str | os.PathLike | typing.BinaryIO | typing.IO[bytes]
model: torch.nn.Module
optimizer: torch.optim.Optimizer
```

Implement the `[adapters.run_save_checkpoint]` and `[adapters.run_load_checkpoint]` adapters, and make sure they pass `uv run pytest -k test_checkpointing`.

### 5.3 Training loop

Now, it's finally time to put all of the components you implemented together into your main training script. It will pay off to make it easy to start training runs with different hyperparameters (e.g., by taking them as command-line arguments), since you will be doing these many times later to study how different choices impact training.

---

**Problem (training\_together): Put it together (4 points)**

---

**Deliverable:** Write a script that runs a training loop to train your model on user-provided input. In particular, we recommend that your training script allow for (at least) the following:

- Ability to configure and control the various model and optimizer hyperparameters.
- Memory-efficient loading of training and validation large datasets with `np.memmap`.
- Serializing checkpoints to a user-provided path.
- Periodically logging training and validation performance (e.g., to console and/or an external service like Weights and Biases).<sup>a</sup>

---

<sup>a</sup> [wandb.ai](https://wandb.ai)

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snyder/cs336\_basics/checkpointing.py

```
1 from typing import IO, BinaryIO
2 import torch
3 import os
4
5
6 def save_checkpoint(
7     model: torch.nn.Module,
8     optimizers: list[torch.optim.Optimizer] | torch.optim.Optimizer,
9     iteration: int,
10    out: str | os.PathLike | BinaryIO | IO[bytes],
11 ):
12     # Extract the original model from a compiled module if present
13     orig_model = model._orig_mod if hasattr(model, "_orig_mod") else model
14
15     if isinstance(optimizers, torch.optim.Optimizer):
16         optimizers = [optimizers]
17
18     torch.save(
19         {
20             "model": orig_model.state_dict(),
21             "optimizer": [optimizer.state_dict() for optimizer in optimizers],
22             "iteration": iteration,
23         },
24         out,
25     )
26
27
28 def load_checkpoint(
29     src: str | os.PathLike | BinaryIO | IO[bytes],
30     model: torch.nn.Module | None = None,
31     optimizers: list[torch.optim.Optimizer] | torch.optim.Optimizer | None = None,
32 ):
33     checkpoint = torch.load(src)
34
35     if model is not None:
36         model.load_state_dict(checkpoint["model"])
37
38     if optimizers is not None:
39         if isinstance(optimizers, torch.optim.Optimizer):
40             optimizers = [optimizers]
41
42         for optimizer, state_dict in zip(optimizers, checkpoint["optimizer"]):
43             optimizer.load_state_dict(state_dict)
44
45     return checkpoint["iteration"]
46
```

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snider/cs336\_basics/train.py

```

1 import torch
2 import os
3 import time
4 import numpy as np
5 import argparse
6 import json
7 import wandb.wandb_run
8 import yaml
9 import math
10 import wandb
11
12 from cs336_basics.adamw import AdamW
13 from cs336_basics.checkpointing import save_checkpoint, load_checkpoint
14 from cs336_basics.lr_schedule import lr_linear_schedule
15 from cs336_basics.model import Transformer
16 from cs336_basics.data_loader import get_batch
17 from cs336_basics.loss import cross_entropy_loss
18 from cs336_basics.gradient_clip import gradient_clip
19 from cs336_basics.lr_schedule import lr_cosine_schedule, lr_double_schedule
20
21
22 class Logger:
23     """Logger that handles console, file, and wandb logging"""
24
25     def __init__(self, log_file: str | None = None, wandb_run: wandb.wandb_run.Run |
26 None = None, resume: bool = False):
27         self.log_file = log_file
28         self.wandb_run = wandb_run
29
30         # Initialize log file
31         if self.log_file:
32             os.makedirs(os.path.dirname(self.log_file), exist_ok=True)
33             # Only clear the file if not resuming from checkpoint
34             if not resume:
35                 with open(self.log_file, "w") as _: # Clear the file
36                     pass
37
38     def log_info(self, message: str | dict, console=True):
39         """Log a message to console and/or file"""
40         if isinstance(message, dict):
41             message = self.format_metrics(message)
42
43         if console:
44             print(message)
45
46         if self.log_file:
47             with open(self.log_file, "a") as f:
48                 f.write(message + "\n")

```

```

48
49 def log_metrics(self, metrics: dict):
50     """Log metrics to wandb"""
51     if self.wandb_run:
52         self.wandb_run.log(metrics)
53
54 def format_metrics(self, metrics_dict: dict) -> str:
55     """Format metrics dictionary into a readable string"""
56     return " | ".join(f"{key}: {value}" for key, value in metrics_dict.items())
57
58
59 class Config(dict):
60     """Config object that allows attribute-style access to dictionary keys
61
62     e.g. config = Config({"training": {"lr": 0.001, "batch_size": 64}})
63
64     can be called as config["training"]["lr"] or config.training.lr
65
66     """
67
68     def __init__(self, *args, **kwargs):
69         super().__init__(*args, **kwargs)
70         self.__dict__ = self
71
72         # Convert nested dictionaries to Config objects
73         for key, value in self.items():
74             if isinstance(value, dict):
75                 self[key] = Config(value)
76
77
78 def load_config_from_file(config_path: str) -> dict:
79     """Load configuration from a file and return as a dictionary"""
80     ext = os.path.splitext(config_path)[1].lower()
81     with open(config_path) as f:
82         if ext == ".json":
83             return json.load(f)
84         elif ext in [".yaml", ".yml"]: # yaml is a nested dictionary
85             return yaml.safe_load(f)
86         else:
87             raise ValueError(f"Unsupported config file format: {ext}")
88
89
90 def load_config(config_path: str | None = None, base_config: dict | None = None) -> Config:
91     """Load configuration from a file and detect runtime device"""
92     # Start with base_config if provided (e.g. when resuming), otherwise load default
93     if base_config is None:
94         default_config_path = os.path.join(os.path.dirname(__file__),
95     ".configs/default.yml")
96         config = load_config_from_file(default_config_path)
97     else:

```

```

97         config = base_config
98
99     # If user specified a config, override defaults
100     if config_path:
101         user_config = load_config_from_file(config_path)
102
103     # Deep update the config
104     for section, section_config in user_config.items():
105         if section in config:
106             config[section].update(section_config)
107         else:
108             config[section] = section_config
109
110     config["run"]["run_id"] = config["run"]["run_id"].replace("<timestamp>", f"
{int(time.time())}")
111
112     # Detect device and dtype at runtime
113     device = "cpu"
114     if torch.cuda.is_available():
115         if config["training"].get("device", None) is not None:
116             device = config["training"]["device"]
117         else:
118             device = "cuda"
119     elif hasattr(torch.backends, "mps") and torch.backends.mps.is_available():
120         device = "mps"
121
122     print(f"Using device: {device}")
123
124     dtype = torch.bfloat16 if torch.cuda.is_available() else torch.float32
125
126     config["device"] = device
127     config["dtype"] = str(dtype) # Convert dtype to string for JSON serialization
128
129     # Convert nested dictionary to Config object
130     return Config(config)
131
132
133 def train(config: Config | None = None):
134     """Train a transformer model with the given configuration"""
135     if config is None:
136         config = load_config()
137
138     # Check if we're resuming from a checkpoint
139     resuming = config.training.get("resume", False)
140     start_step = 1
141
142     run_dir = os.path.join(config.run.out_dir, config.run.run_id)
143     config_outfile = os.path.join(run_dir, "config.json")
144     log_file = os.path.join(run_dir, "log.txt")
145     checkpoint_dir = os.path.join(run_dir, "checkpoints")

```

```

146
147 os.makedirs(run_dir, exist_ok=True)
148 os.makedirs(checkpoint_dir, exist_ok=True)
149
150 # Create symlink pointing `latest` to run_dir (remove `latest` if it exists)
151 latest_symlink = os.path.join(config.run.out_dir, "latest")
152 if os.path.islink(latest_symlink) or os.path.exists(latest_symlink):
153     os.remove(latest_symlink)
154 os.symlink(os.path.abspath(run_dir), latest_symlink, target_is_directory=True)
155
156 # Initialize wandb and logger
157 wandb_run = (
158     None
159     if not config.run.wandb_project
160     else wandb.init(
161         project=config.run.wandb_project,
162         id=config.run.run_id, # Use run_id as the wandb id
163         resume="must" if resuming else None, # Set resume conditionally
164         name=config.run.run_id,
165         config=config,
166         dir=run_dir,
167         tags=config.run.wandb_tags,
168     )
169 )
170
171 logger = Logger(log_file=log_file, wandb_run=wandb_run, resume=resuming)
172
173 # Save configuration (only if not resuming)
174 if resuming:
175     logger.log_info(f"Resuming training from existing config: {config_outfile}")
176 else:
177     with open(config_outfile, "w") as f:
178         json.dump(config, f, indent=2, default=lambda o: list(o) if isinstance(o,
tuple) else o.__dict__)
179     logger.log_info(f"Saved config to: {config_outfile}")
180
181 device = config.device
182 dtype = getattr(torch, config.dtype.split(".")[1]) # Convert string back to
torch dtype
183
184 train_data = np.memmap(config.data.train_data_path, dtype=np.uint16, mode="r")
185 valid_data = np.memmap(config.data.valid_data_path, dtype=np.uint16, mode="r")
186
187 # Initialize model
188 model = Transformer(**config.model, device=device, dtype=dtype)
189 model.to(device)
190
191 print(f"Trainable params: {sum(p.numel() for p in model.parameters() if
p.requires_grad)}")
192
193 # Only decay 2D parameters (i.e. not layernorms)

```

```

194 param_dict = {pn: p for pn, p in model.named_parameters() if p.requires_grad}
195 decay_params = [p for n, p in param_dict.items() if p.dim() >= 2]
196 nodecay_params = [p for n, p in param_dict.items() if p.dim() < 2]
197 optim_groups = [
198     {"params": decay_params, **config.optimizer},
199     {"params": nodecay_params, **config.optimizer, "weight_decay": 0.0},
200 ]
201 num_decay_params = sum(p.numel() for p in decay_params)
202 num_nodecay_params = sum(p.numel() for p in nodecay_params)
203 print(f"Decayed parameter tensors: {len(decay_params)}, with {num_decay_params:,}
parameters")
204 print(f"Non-decayed parameter tensors: {len(nodecay_params)}, with
{num_nodecay_params:,} parameters")
205
206 optimizer = AdamW(optim_groups, **config.optimizer)
207 # optimizer = AdamW(model.parameters(), **config.optimizer)
208
209 # Load checkpoint if resuming
210 if resuming:
211     checkpoint_path = config.training.resume_checkpoint
212     logger.log_info(f"Loading checkpoint from: {checkpoint_path}")
213     start_step = load_checkpoint(checkpoint_path, model, optimizer) + 1
214     logger.log_info(f"Resuming training from step {start_step}")
215
216 # Compile + AMP on GPU, AOT on MPS
217 use_compile = True
218 if use_compile and device != "mps":
219     model = torch.compile(model)
220     torch.set_float32_matmul_precision("high")
221 elif use_compile and device == "mps":
222     model = torch.compile(model, backend="aot_eager") # this does not accelerate,
but does check for errors via compilation before running
223
224 max_steps = config.training.max_steps
225 batch_size = config.training.batch_size
226 max_l2_norm = config.training.max_l2_norm
227 eval_interval = config.training.eval_interval
228 checkpoint_interval = config.training.checkpoint_interval
229 grad_accum_steps = config.training.grad_accum_steps
230
231 lr_max = config.training.lr_max
232 lr_inter = config.training.lr_inter
233 lr_min = config.training.lr_min
234 warmup_ratio = config.training.warmup_ratio
235 warmup_iters = config.training.warmup_iters
236 phase_one_iters = config.training.phase_one_iters
237 phase_two_iters = config.training.phase_two_iters
238 phase_two_type = config.training.phase_two_type
239 cosine_cycle_iters = config.training.cosine_cycle_iters
240 linear_cycle_iters = config.training.linear_cycle_iters
241

```



```
242     if warmup_iters is False or warmup_iters is None:
243         warmup_iters = int(warmup_ratio * max_steps)
244
245     if cosine_cycle_iters is False or cosine_cycle_iters is None:
246         cosine_cycle_iters = max_steps
247
248     if linear_cycle_iters is False or linear_cycle_iters is None:
249         linear_cycle_iters = max_steps
250
251     if phase_two_iters is False or phase_two_iters is None:
252         phase_two_iters = max_steps
253
254     def evaluate(step: int, is_last_step: bool):
255         n_eval_steps = config.training.eval_steps
256
257         if is_last_step:
258             n_eval_steps = n_eval_steps * 3
259
260         model.eval()
261
262         with torch.no_grad():
263             val_loss = 0.0
264             for _ in range(n_eval_steps):
265                 x, y = get_batch(valid_data, config.training.eval_batch_size,
config.model.context_length, device)
266                 with torch.autocast(device_type=device, dtype=dtype):
267                     logits = model(x)
268                     loss = cross_entropy_loss(logits, y)
269                     val_loss += loss.item()
270             val_loss /= n_eval_steps
271
272         progress_str = get_progress_str(step, max_steps)
273
274         # WandB metrics
275         metrics = {
276             "eval/loss": val_loss,
277             "eval/perplexity": get_perplexity(val_loss),
278             "eval/peak_memory": get_peak_memory(device),
279             "step": step,
280         }
281
282         # Console + local file metrics
283         display_metrics = {
284             "step": progress_str,
285             "v_loss": f"{val_loss:.4f}",
286             "v_ppl": f"{get_perplexity(val_loss):.2f}",
287             "mem": f"{get_peak_memory(device):.1f}MB",
288         }
289
290         logger.log_info(display_metrics)
```

```

291         logger.log_metrics(metrics)
292
293         # Restore to training mode
294         model.train()
295
296         # Only evaluate before training if not resuming
297         if config.training.eval_before_training and not resuming:
298             evaluate(0)
299
300         # Load the first batch
301         x, y = get_batch(train_data, batch_size, config.model.context_length, device)
302
303         model.train()
304
305         for step in range(start_step, max_steps + 1):
306             t0 = time.time()
307             is_last_step = step == max_steps
308             loss_accum = 0.0
309
310             # Gradient accumulation loop
311             for _ in range(grad_accum_steps):
312                 with torch.autocast(device_type=device, dtype=dtype):
313                     logits = model(x)
314
315                     loss = cross_entropy_loss(logits, y) / grad_accum_steps
316
317             x, y = get_batch(train_data, batch_size, config.model.context_length,
device)
318
319             loss_accum += loss.detach()
320             loss.backward()
321
322             norm = gradient_clip(model.parameters(), max_l2_norm)
323
324             if config.training.lr_schedule == "linear":
325                 lr = lr_linear_schedule(step, lr_max, lr_min, warmup_iters,
linear_cycle_iters)
326             elif config.training.lr_schedule == "cosine":
327                 lr = lr_cosine_schedule(step, lr_max, lr_min, warmup_iters,
cosine_cycle_iters)
328             elif config.training.lr_schedule == "double":
329                 lr = lr_double_schedule(
330                     step, lr_max, lr_inter, lr_min, warmup_iters, phase_one_iters,
phase_two_iters, phase_two_type
331                 )
332
333             for param_group in optimizer.param_groups:
334                 param_group["lr"] = lr
335
336             optimizer.step()
337             optimizer.zero_grad(set_to_none=True)

```

```

338
339     if device == "cuda":
340         torch.cuda.synchronize()
341     elif device == "mps":
342         torch.mps.synchronize()
343
344     t1 = time.time()
345     dt = t1 - t0
346     tokens_per_sec = config.model.context_length * batch_size * grad_accum_steps
    / dt
347     train_loss = loss_accum.item()
348     progress_str = get_progress_str(step, max_steps)
349
350     # WandB metrics
351     metrics = {
352         "train/loss": train_loss,
353         "train/perplexity": get_perplexity(train_loss),
354         "train/lr": lr,
355         "train/grad_norm": norm,
356         "train/tokens_per_sec": tokens_per_sec,
357         "train/peak_memory": get_peak_memory(device),
358         "step": step,
359     }
360
361     # Console + local file metrics
362     display_metrics = {
363         "step": progress_str,
364         "t_loss": f"{train_loss:.4f}",
365         "t_ppl": f"{get_perplexity(train_loss):.2f}",
366         "lr": f"{lr:.4e}",
367         "grad_norm": f"{norm:.2f}",
368         "mem": f"{get_peak_memory(device):.1f}MB",
369         "tok/sec": f"{int(tokens_per_sec):,}",
370         "dt": f"{dt * 1000:.2f}ms",
371     }
372
373     logger.log_info(display_metrics)
374     logger.log_metrics(metrics)
375
376     if step % eval_interval == 0 or is_last_step:
377         evaluate(step, is_last_step)
378
379     if step % checkpoint_interval == 0 or is_last_step:
380         checkpoint_path = os.path.join(checkpoint_dir, f"checkpoint_{step}.pt")
381         save_checkpoint(model, optimizer, step, checkpoint_path)
382
383         # Create symlink pointing `latest` to checkpoint_path (remove `latest` if
    it exists)
384         latest_symlink = os.path.join(checkpoint_dir, "latest.pt")
385         if os.path.islink(latest_symlink) or os.path.exists(latest_symlink):
386             os.remove(latest_symlink)

```

```
387         os.symlink(os.path.abspath(checkpoint_path), latest_symlink)
388
389         logger.log_info(f"Saved checkpoint to: {checkpoint_path}")
390
391     wandb.finish()
392
393
394 def get_peak_memory(device):
395     """Get peak memory usage in MB on the current device"""
396     if device != "cuda":
397         return 0
398
399     peak_memory = torch.cuda.max_memory_allocated() / (1024 * 1024)
400     torch.cuda.reset_peak_memory_stats()
401     return peak_memory
402
403
404 def get_perplexity(loss):
405     """Calculate perplexity from cross-entropy loss"""
406     return math.exp(min(loss, 20)) # Cap at 20 to avoid overflow
407
408
409 def get_progress_str(step, max_steps):
410     return f"step {step}/{max_steps} ({step / max_steps * 100:.2f}%)"
411
412
413 def parse_value(value_str: str):
414     """Convert argparse arg to list, int, float, bool, or leave as string."""
415     if value_str.strip().startswith("[") and value_str.strip().endswith("]"):
416         content = value_str.strip()[1:-1].strip()
417         return [parse_value(v.strip()) for v in content.split(",")] if content else
418     []
419
420     try:
421         return int(value_str)
422     except ValueError:
423         pass
424
425     try:
426         return float(value_str)
427     except ValueError:
428         pass
429
430     lower = value_str.strip().lower()
431     if lower in ("true", "false"):
432         return lower == "true"
433
434     return value_str
435
```

```

436 def deep_set(config_dict, key_path: str, value):
437     """Deeply set dot-separated key in a config dictionary"""
438     keys = key_path.split(".")
439     d = config_dict
440     for k in keys[:-1]:
441         if k not in d or not isinstance(d[k], dict):
442             d[k] = {}
443         d = d[k]
444     d[keys[-1]] = value
445
446
447 if __name__ == "__main__":
448     parser = argparse.ArgumentParser(description="Train a transformer model")
449     parser.add_argument("--config", type=str, help="Path to config file (optional)")
450     parser.add_argument("--resume-from", type=str, help="Path to run directory to
resume from")
451     parser.add_argument("--override-config", type=str, help="Path to config file with
values to override when resuming")
452     parser.add_argument(
453         "--override-param",
454         action="append",
455         default=[],
456         help="Override a config param, e.g. model.d_model=512 (can be repeated)",
457     )
458     args = parser.parse_args()
459
460     if args.resume_from:
461         # Load config from the previous run
462         resume_config_path = os.path.join(args.resume_from, "config.json")
463         base_config = load_config_from_file(resume_config_path)
464         base_config["training"]["resume"] = True
465         base_config["training"]["resume_checkpoint"] = os.path.join(args.resume_from,
"checkpoints/latest.pt")
466
467         # Use the override config if provided, or None otherwise
468         config = load_config(args.override_config, base_config=base_config)
469     else:
470         config = load_config(args.config) # Will use default if args.config is None
471
472     for override_str in args.override_param:
473         if "=" not in override_str:
474             raise ValueError(f"Invalid override: {override_str}, must be like
key=val")
475
476         key, raw_value = override_str.split("=", 1)
477         value = parse_value(raw_value)
478         deep_set(config, key, value)
479
480     train(config)
481

```

## 6 Generating text

Now that we can train models, the last piece we need is the ability to generate text from our model. Recall that a language model takes in a (possibly batched) integer sequence of length (`sequence_length`) and produces a matrix of size (`sequence_length × vocab_size`), where each element of the sequence is a probability distribution predicting the next word after that position. We will now write a few functions to turn this into a sampling scheme for new sequences.

**Softmax** By standard convention, the language model output is the output of the final linear layer (the “logits”) and so we have to turn this into a normalized probability via the *softmax* operation, which we saw earlier in Eq 10.

**Decoding** To generate text (decode) from our model, we will provide the model with a sequence of prefix tokens (the “prompt”), and ask it to produce a probability distribution over the vocabulary that predicts the next word in the sequence. Then, we will sample from this distribution over the vocabulary items to determine the next output token.

Concretely, one step of the decoding process should take in a sequence  $x_{1...t}$  and return a token  $x_{t+1}$  via the following equation,

$$P(x_{t+1} = i \mid x_{1...t}) = \frac{\exp(v_i)}{\sum_j \exp(v_j)}$$

$$v = \text{TransformerLM}(x_{1...t})_t \in \mathbb{R}^{\text{vocab\_size}}$$

where `TransformerLM` is our model which takes as input a sequence of `sequence_length` and produces a matrix of size (`sequence_length × vocab_size`), and we take the last element of this matrix, as we are looking for the next word prediction at the  $t$ -th position.

This gives us a basic decoder by repeatedly sampling from these one-step conditionals (appending our previously-generated output token to the input of the next decoding timestep) until we generate the end-of-sequence token `<|endoftext|>` (or a user-specified maximum number of tokens to generate).

**Decoder tricks** We will be experimenting with small models, and small models can sometimes generate very low quality texts. Two simple decoder tricks can help fix these issues. First, in *temperature scaling* we modify our softmax with a temperature parameter  $\tau$ , where the new softmax is

$$\text{softmax}(v, \tau)_i = \frac{\exp(v_i/\tau)}{\sum_{j=1}^{\text{vocab\_size}} \exp(v_j/\tau)}. \quad (24)$$

Note how setting  $\tau \rightarrow 0$  makes it so that the largest element of  $v$  dominates, and the output of the softmax becomes a one-hot vector concentrated at this maximal element.

Second, another trick is *nucleus* or *top-p* sampling, where we modify the sampling distribution by truncating low-probability words. Let  $q$  be a probability distribution that we get from a (temperature-scaled) softmax of size (`vocab_size`). Nucleus sampling with hyperparameter  $p$  produces the next token according to the equation

$$P(x_{t+1} = i \mid q) = \begin{cases} \frac{q_i}{\sum_{j \in V(p)} q_j} & \text{if } i \in V(p) \\ 0 & \text{otherwise} \end{cases}$$

where  $V(p)$  is the *smallest* set of indices such that  $\sum_{j \in V(p)} q_j \geq p$ . You can compute this quantity easily by first sorting the probability distribution  $q$  by magnitude, and selecting the largest vocabulary elements until you reach the target level of  $\alpha$ .

### Problem (decoding): Decoding (3 points)

---

**Deliverable:** Implement a function to decode from your language model. We recommend that you support the following features:

- Generate completions for a user-provided prompt (i.e., take in some  $x_{1...t}$  and sample a completion until you hit an `<|endoftext|>` token).
- Allow the user to control the maximum number of generated tokens.
- Given a desired temperature value, apply softmax temperature scaling to the predicted next-word distributions before sampling.
- Top-p sampling (Holtzman et al., 2020; also referred to as nucleus sampling), given a user-specified threshold value.

~/Dropbox/Eric/Machine Learning/cs336/cs336-a1-main-brandon-snyder/cs336\_basics/decoding.py

```
1 import torch
2 from cs336_basics.tokenizer import Tokenizer
3 from cs336_basics.model import Transformer
4
5
6 def decode(
7     model: Transformer,
8     tokenizer: Tokenizer,
9     prompt: str,
10    max_new_tokens: int = 32,
11    temperature: float = 0.7,
12    top_p: float = 0.9,
13 ):
14     end_id = tokenizer.encode("<|endoftext|>")[0]
15     input_ids = tokenizer.encode(prompt)
16     device = next(model.parameters()).device
17     context_length = model.context_length
18
19     with torch.no_grad():
20         for _ in range(max_new_tokens):
21             window_input_ids = input_ids[-context_length:] if len(input_ids) >=
context_length else input_ids
22             x = torch.tensor([window_input_ids], dtype=torch.long, device=device)
23
24             logits = model(x)
25             next_logits = logits[0, -1, :]
26
27             scaled = next_logits / temperature
28             stable = scaled - scaled.max()
29             exp_vals = stable.exp()
30             probs = exp_vals / exp_vals.sum()
31
32             sorted_probs, sorted_idxs = torch.sort(probs, descending=True)
33             cumsum = torch.cumsum(sorted_probs, dim=0)
34             cutoff_idx = torch.searchsorted(cumsum, top_p)
35             trimmed_probs = sorted_probs[: cutoff_idx + 1]
36             trimmed_idxs = sorted_idxs[: cutoff_idx + 1]
37             trimmed_probs /= trimmed_probs.sum()
38
39             next_token = trimmed_idxs[torch.multinomial(trimmed_probs, 1).item()]
40             if next_token.item() == end_id:
41                 break
42             input_ids.append(next_token.item())
43
44     return tokenizer.decode(input_ids)
45
```



## 7 Experiments

Now it is time to put everything together and train (small) language models on a pretraining dataset.

### 7.1 How to Run Experiments and Deliverables

The best way to understand the rationale behind the architectural components of a Transformer is to actually modify it and run it yourself. There is no substitute for hands-on experience.

To this end, it's important to be able to experiment **quickly, consistently, and keep records** of what you did. To experiment quickly, we will be running many experiments on a small scale model (17M parameters) and simple dataset (TinyStories). To do things consistently, you will ablate components and vary hyperparameters in a systematic way, and to keep records we will ask you to submit a log of your experiments and learning curves associated with each experiment.

To make it possible to submit loss curves, **make sure to periodically evaluate validation losses and record both the number of steps and wallclock times**. You might find logging infrastructure such as Weights and Biases helpful.

#### Problem (experiment\_log): Experiment logging (3 points)

For your training and evaluation code, create experiment tracking infrastructure that allows you to track your experiments and loss curves with respect to gradient steps and wallclock time.

**Deliverable:** Logging infrastructure code for your experiments and an experiment log (a document of all the things you tried) for the assignment problems below in this section.

### 7.2 TinyStories

We are going to start with a very simple dataset (TinyStories; Eldan and Li, 2023) where models will train quickly, and we can see some interesting behaviors. The instructions for getting this dataset is at section 1. An example of what this dataset looks like is below.

#### Example (tinystories\_example): One example from TinyStories

Once upon a time there was a little boy named Ben. Ben loved to explore the world around him. He saw many amazing things, like beautiful vases that were on display in a store. One day, Ben was walking through the store when he came across a very special vase. When Ben saw it he was amazed! He said, "Wow, that is a really amazing vase! Can I buy it?" The shopkeeper smiled and said, "Of course you can. You can take it home and show all your friends how amazing it is!" So Ben took the vase home and he was so proud of it! He called his friends over and showed them the amazing vase. All his friends thought the vase was beautiful and couldn't believe how lucky Ben was. And that's how Ben found an amazing vase in the store!

**Hyperparameter tuning** We will tell you some very basic hyperparameters to start with and ask you to find some settings for others that work well.

**vocab\_size** 10000. Typical vocabulary sizes are in the tens to hundreds of thousands. You should vary this and see how the vocabulary and model behavior changes.

**context\_length** 256. Simple datasets such as TinyStories might not need long sequence lengths, but for the later OpenWebText data, you may want to vary this. Try varying this and seeing the impact on both the per-iteration runtime and the final perplexity.

**d\_model** 512. This is slightly smaller than the 768 dimensions used in many small Transformer papers, but this will make things faster.

**d\_ff** 1344. This is roughly  $\frac{8}{3}d_{\text{model}}$  while being a multiple of 64, which is good for GPU performance.

**RoPE theta parameter**  $\Theta$  10000.

**number of layers and heads** 4 layers, 16 heads. Together, this will give about 17M non-embedding parameters which is a fairly small Transformer.

**total tokens processed** 327,680,000 (your batch size  $\times$  total step count  $\times$  context length should equal roughly this value).

You should do some trial and error to find good defaults for the following other hyperparameters: **learning rate**, **learning rate warmup**, **other AdamW hyperparameters** ( $\beta_1, \beta_2, \epsilon$ ), and **weight decay**. You can find some typical choices of such hyperparameters in [Kingma and Ba \[2015\]](#).

**Putting it together** Now you can put everything together by getting a trained BPE tokenizer, tokenizing the training dataset, and running this in the training loop that you wrote. **Important note:** If your implementation is correct and efficient, the above hyperparameters should result in a roughly 30-40 minute runtime on 1 H100 GPU. If you have runtimes that are much longer, please check and make sure your dataloading, checkpointing, or validation loss code is not bottlenecking your runtimes and that your implementation is properly batched.

**Tips and tricks for debugging model architectures** We highly recommend getting comfortable with your IDE's built-in debugger (e.g., VSCode/PyCharm), which will save you time compared to debugging with print statements. If you use a text editor, you can use something more like `pdb`. A few other good practices when debugging model architectures are:

- A common first step when developing any neural net architecture is to overfit to a single minibatch. If your implementation is correct, you should be able to quickly drive the training loss to near-zero.
- Set debug breakpoints in various model components, and inspect the shapes of intermediate tensors to make sure they match your expectations.
- Monitor the norms of activations, model weights, and gradients to make sure they are not exploding or vanishing.

#### **Problem (learning\_rate): Tune the learning rate (3 points) (4 H100 hrs)**

The learning rate is one of the most important hyperparameters to tune. Taking the base model you've trained, answer the following questions:

- (a) Perform a hyperparameter sweep over the learning rates and report the final losses (or note divergence if the optimizer diverges).

**Deliverable:** Learning curves associated with multiple learning rates. Explain your hyperparameter search strategy.

**Deliverable:** A model with validation loss (per-token) on TinyStories of at most 1.45

### Low-Resource/Downscaling Tip: Train for few steps on CPU or Apple Silicon

If you are running on `cpu` or `mps`, you should instead reduce the total tokens processed count to 40,000,000, which will be sufficient to produce reasonably fluent text. You may also increase the target validation loss from 1.45 to 2.00.

Running our solution code with a tuned learning rate on an M3 Max chip and 36 GB of RAM, we use  $\text{batch size} \times \text{total step count} \times \text{context length} = 32 \times 5000 \times 256 = 40,960,000$  tokens, which takes 1 hour and 22 minutes on `cpu` and 36 minutes on `mps`. At step 5000, we achieve a validation loss of 1.80.

Some additional tips:

- When using  $X$  training steps, we suggest adjusting the cosine learning rate decay schedule to terminate its decay (i.e., reach the minimum learning rate) at precisely step  $X$ .

- When using `mps`, do **not** use TF32 kernels, i.e., do **not** set

```
torch.set_float32_matmul_precision('high')
```

as you might with `cuda` devices. We tried enabling TF32 kernels with `mps` (torch version 2.6.0) and found the backend will use silently broken kernels that cause unstable training.

- You can speed up training by JIT-compiling your model with `torch.compile`. Specifically:

- On `cpu`, compile your model with

```
model = torch.compile(model)
```

- On `mps`, you can somewhat optimize the backward pass using

```
model = torch.compile(model, backend="aot_eager")
```

Compilation with Inductor is not supported on `mps` as of torch version 2.6.0.

- (b) Folk wisdom is that the best learning rate is “at the edge of stability.” Investigate how the point at which learning rates diverge is related to your best learning rate.

**Deliverable:** Learning curves of increasing learning rate which include at least one divergent run and an analysis of how this relates to convergence rates.

Now let’s vary the batch size and see what happens to training. Batch sizes are important – they let us get higher efficiency from our GPUs by doing larger matrix multiplies, but is it true that we always want batch sizes to be large? Let’s run some experiments to find out.

### **Problem (batch\_size\_experiment): Batch size variations (1 point) (2 H100 hrs)**

Vary your batch size all the way from 1 to the GPU memory limit. Try at least a few batch sizes in between, including typical sizes like 64 and 128.

**Deliverable:** Learning curves for runs with different batch sizes. The learning rates should be optimized again if necessary.

**Deliverable:** A few sentences discussing of your findings on batch sizes and their impacts on training.

With your decoder in hand, we can now generate text! We will generate from the model and see how good it is. As a reference, you should get outputs that look at least as good as the example below.

**Example (ts\_generate\_example): Sample output from a TinyStories language model**

Once upon a time, there was a pretty girl named Lily. She loved to eat gum, especially the big black one. One day, Lily's mom asked her to help cook dinner. Lily was so excited! She loved to help her mom. Lily's mom made a big pot of soup for dinner. Lily was so happy and said, "Thank you, Mommy! I love you." She helped her mom pour the soup into a big bowl. After dinner, Lily's mom made some yummy soup. Lily loved it! She said, "Thank you, Mommy! This soup is so yummy!" Her mom smiled and said, "I'm glad you like it, Lily." They finished cooking and continued to cook together. The end.

**Low-Resource/Downscaling Tip: Generate text on CPU or Apple Silicon**

If instead you used the low-resource configuration with 40M tokens processed, you should see generations that still resemble English but are not as fluent as above. For example, our sample output from a TinyStories language model trained on 40M tokens is below:

Once upon a time, there was a little girl named Sue. Sue had a tooth that she loved very much. It was his best head. One day, Sue went for a walk and met a ladybug! They became good friends and played on the path together.

"Hey, Polly! Let's go out!" said Tim. Sue looked at the sky and saw that it was difficult to find a way to dance shining. She smiled and agreed to help the talking!"

As Sue watched the sky moved, what it was. She

Here is the precise problem statement and what we ask for:

**Problem (generate): Generate text (1 point)**

Using your decoder and your trained checkpoint, report the text generated by your model. You may need to manipulate decoder parameters (temperature, top-p, etc.) to get fluent outputs.

**Deliverable:** Text dump of at least 256 tokens of text (or until the first `<|endoftext|>` token), and a brief comment on the fluency of this output and at least two factors which affect how good or bad this output is.

### 7.3 Ablations and architecture modification

The best way to understand the Transformer is to actually modify it and see how it behaves. We will now do a few simple ablations and modifications.

**Ablation 1: layer normalization** It is often said that layer normalization is important for the stability of Transformer training. But perhaps we want to live dangerously. Let's remove RMSNorm from each of our Transformer blocks and see what happens.

**Problem (layer\_norm\_ablation): Remove RMSNorm and train (1 point) (1 H100 hr)**

Remove all of the RMSNorms from your Transformer and train. What happens at the previous optimal learning rate? Can you get stability by using a lower learning rate?

**Deliverable:** A learning curve for when you remove RMSNorms and train, as well as a learning curve for the best learning rate.

**Deliverable:** A few sentence commentary on the impact of RMSNorm.

Let's now investigate another layer normalization choice that seems arbitrary at first glance. *Pre-norm* Transformer blocks are defined as

$$\begin{aligned} z &= x + \text{MultiHeadedSelfAttention}(\text{RMSNorm}(x)) \\ y &= z + \text{FFN}(\text{RMSNorm}(z)). \end{aligned}$$

This is one of the few ‘consensus’ modifications to the original Transformer architecture, which used a *post-norm* approach as

$$\begin{aligned} z &= \text{RMSNorm}(x + \text{MultiHeadedSelfAttention}(x)) \\ y &= \text{RMSNorm}(z + \text{FFN}(z)). \end{aligned}$$

Let's revert back to the *post-norm* approach and see what happens.

**Problem (pre\_norm\_ablation): Implement post-norm and train (1 point) (1 H100 hr)**

Modify your pre-norm Transformer implementation into a post-norm one. Train with the post-norm model and see what happens.

**Deliverable:** A learning curve for a post-norm transformer, compared to the pre-norm one.

We see that layer normalization has a major impact on the behavior of the transformer, and that even the position of the layer normalization is important.

**Ablation 2: position embeddings** We will next investigate the impact of the position embeddings on the performance of the model. Specifically, we will compare our base model (with RoPE) with not including position embeddings at all (NoPE). It turns out that decoder-only transformers, i.e., those with a causal mask as we have implemented, can in theory *infer relative or absolute position information* without being provided with position embeddings explicitly [Tsai et al., 2019, Kazemnejad et al., 2023]. We will now test empirically how NoPE performs compare to RoPE.

**Problem (no\_pos\_emb): Implement NoPE (1 point) (1 H100 hr)**

Modify your Transformer implementation with RoPE to remove the position embedding information entirely, and see what happens.

**Deliverable:** A learning curve comparing the performance of RoPE and NoPE.

**Ablation 3: SwiGLU vs. SiLU** Next, we will follow Shazeer [2020] and test the importance of gating in the feed-forward network, by comparing the performance of SwiGLU feed-forward networks versus feed-forward networks using SiLU activations but no gated linear unit (GLU):

$$\text{FFN}_{\text{SiLU}}(x) = W_2 \text{SiLU}(W_1 x). \quad (25)$$

Recall that in our SwiGLU implementation, we set the dimensionality of the inner feed-forward layer to be roughly  $d_{\text{ff}} = \frac{8}{3}d_{\text{model}}$  (while ensuring that  $d_{\text{ff}} \bmod 64 = 0$ , to make use of GPU tensor cores). In your  $\text{FFN}_{\text{SiLU}}$  implementation you should set  $d_{\text{ff}} = 4 \times d_{\text{model}}$ , to approximately match the parameter count of the SwiGLU feed-forward network (which has three instead of two weight matrices).

**Problem (swiglu\_ablation): SwiGLU vs. SiLU (1 point) (1 H100 hr)**

**Deliverable:** A learning curve comparing the performance of SwiGLU and SiLU feed-forward networks, with approximately matched parameter counts.

**Deliverable:** A few sentences discussing your findings.

**Low-Resource/Downscaling Tip: Online students with limited GPU resources should test modifications on TinyStories**

In the remainder of the assignment, we will move to a larger-scale, noisier web dataset (OpenWebText), experimenting with architecture modifications and (optionally) making a submission to the course leaderboard.

It takes a long time to train an LM to fluency on OpenWebText, so we suggest that online students with limited GPU access continue testing modifications on TinyStories (using validation loss as a metric to evaluate performance).

## 7.4 Running on OpenWebText

We will now move to a more standard pretraining dataset created from a webcrawl. A small sample of OpenWebText (Gokaslan et al., 2019) is also provided as a single text file: see section 11 for how to access this file.

Here is an example from OpenWebText. Note how the text is much more realistic, complex, and varied. You may want to look through the training dataset to get a sense of what training data looks like for a webscraped corpus.

**Example (owt\_example): One example from OWT**

Baseball Prospectus director of technology Harry Pavlidis took a risk when he hired Jonathan Judge.

Pavlidis knew that, as Alan Schwarz wrote in *The Numbers Game*, “no corner of American culture is more precisely counted, more passionately quantified, than performances of baseball players.” With a few clicks here and there, you can find out that Noah Syndergaard’s fastball revolves more than 2,100 times per minute on its way to the plate, that Nelson Cruz had the game’s highest average exit velocity among qualified hitters in 2016 and myriad other tidbits that seem ripped from a video game or science fiction novel. The rising ocean of data has empowered an increasingly important actor in baseball’s culture: the analytical hobbyist.

That empowerment comes with added scrutiny – on the measurements, but also on the people and publications behind them. With Baseball Prospectus, Pavlidis knew all about the backlash that accompanies quantitative imperfection. He also knew the site’s catching metrics needed to be reworked, and that it would take a learned mind – someone who could tackle complex statistical modeling problems – to complete the job.

“He freaks us out.” Harry Pavlidis

Pavlidis had a hunch that Judge “got it” based on the latter’s writing and their interaction at a site-sponsored ballpark event. Soon thereafter, the two talked over drinks. Pavlidis’ intuition was validated. Judge was a fit for the position – better yet, he was a willing fit. “I spoke to a lot of people,” Pavlidis said, “he was the only one brave enough to take it on.” [...]

**Note:** You may have to re-tune your hyperparameters such as learning rate or batch size for this experiment.

**Problem (main\_experiment): Experiment on OWT (2 points) (3 H100 hrs)**

Train your language model on OpenWebText with the same model architecture and total training iterations as TinyStories. How well does this model do?

**Deliverable:** A learning curve of your language model on OpenWebText. Describe the difference in losses from TinyStories – how should we interpret these losses?

**Deliverable:** Generated text from OpenWebText LM, in the same format as the TinyStories outputs. How is the fluency of this text? Why is the output quality worse even though we have the same model and compute budget as TinyStories?

## 7.5 Your own modification + leaderboard

Congratulations on getting to this point. You're almost done! You will now try to improve upon the Transformer architecture, and see how your hyperparameters and architecture stack up against other students in the class.

**Rules for the leaderboard** There are no restrictions other than the following:

**Runtime** Your submission can run for at most 1.5 hours on an H100. You can enforce this by setting `--time=01:30:00` in your slurm submission script.

**Data** You may only use the OpenWebText training dataset that we provide.

Otherwise, you are free to do whatever your heart desires.

If you are looking for some ideas on what to implement, you can checkout some of these resources:

- State-of-the-art open-source LLM families, such as Llama 3 [Grattafiori et al., 2024] or Qwen 2.5 [Yang et al., 2024].
- The NanoGPT speedrun repository (<https://github.com/KellerJordan/modded-nanogpt>), where community members post many interesting modifications for “speedrunning” small-scale language model pretraining. For example, a common modification that dates back to the original Transformer paper is to tie the weights of the input and output embeddings together (see [Vaswani et al. 2017] (Section 3.4) and [Chowdhery et al. 2022] (Section 2)). If you do try weight tying, you may have to decrease the standard deviation of the embedding/LM head init.

You will want to test these on either a small subset of OpenWebText or on TinyStories before trying the full 1.5-hour run.

As a caveat, we do note that some of the modifications you may find working well in this leaderboard may not generalize to larger-scale pretraining. We will explore this idea further in the scaling laws unit of the course.

### Problem (leaderboard): Leaderboard (6 points) (10 H100 hrs)

You will train a model under the leaderboard rules above with the goal of minimizing the validation loss of your language model within 1.5 H100-hour.

**Deliverable:** The final validation loss that was recorded, an associated learning curve that clearly shows a wallclock-time x-axis that is less than 1.5 hours and a description of what you did. We expect a leaderboard submission to beat at least the naive baseline of a 5.0 loss. Submit to the leaderboard here: <https://github.com/stanford-cs336/assignment1-basics-leaderboard>.

## References

- Ronen Eldan and Yuanzhi Li. TinyStories: How small can language models be and still speak coherent English?, 2023. arXiv:2305.07759.
- Aaron Gokaslan, Vanya Cohen, Ellie Pavlick, and Stefanie Tellex. OpenWebText corpus. <http://Skylion007.github.io/OpenWebTextCorpus>, 2019.
- Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural machine translation of rare words with subword units. In *Proc. of ACL*, 2016.
- Changhan Wang, Kyunghyun Cho, and Jiatao Gu. Neural machine translation with byte-level subwords, 2019. arXiv:1909.03341.
- Philip Gage. A new algorithm for data compression. *C Users Journal*, 12(2):23–38, February 1994. ISSN 0898-9788.
- Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners, 2019.
- Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding by generative pre-training, 2018.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Proc. of NeurIPS*, 2017.
- Toan Q. Nguyen and Julian Salazar. Transformers without tears: Improving the normalization of self-attention. In *Proc. of IWSWLT*, 2019.
- Ruibin Xiong, Yunchang Yang, Di He, Kai Zheng, Shuxin Zheng, Chen Xing, Huishuai Zhang, Yanyan Lan, Liwei Wang, and Tie-Yan Liu. On layer normalization in the Transformer architecture. In *Proc. of ICML*, 2020.
- Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer normalization, 2016. arXiv:1607.06450.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. Llama: Open and efficient foundation language models, 2023. arXiv:2302.13971.
- Biao Zhang and Rico Sennrich. Root mean square layer normalization. In *Proc. of NeurIPS*, 2019.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu



Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani, Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhota, Lauren Rantala-Yearly, Laurens van der Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo, Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang, Olivier Duchenne, Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic, Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu, Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov, Shaoliang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale, Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha, Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gonguet, Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenxin Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide Xia, Xinfeng Xie, Xuchao Jia, Xuwei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei, Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Coudert, Zheng Yan, Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey, Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma, Alex Boesenberg, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo, Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola, Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence, Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu, Changan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich, Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippos Kokkinos, Firat Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi, Zhang, Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha, Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya, Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal, Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A, Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca Wehrstedt, Madian Khabsa, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson, Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally, Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov, Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat, Moham-

- mad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White, Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager, Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang, Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra, Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh, Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Sheng Feng, Shenghao Lin, Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang, Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe, Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou, Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov, Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaojuan Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi, Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu, Wang, Yu Zhao, Yuchen Hao, Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu Yang, Zhiwei Zhao, and Zhiyu Ma. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam Shazeer, Vinodkumar Prabhakaran, Emily Reif, Nan Du, Ben Hutchinson, Reiner Pope, James Bradbury, Jacob Austin, Michael Isard, Guy Gur-Ari, Pengcheng Yin, Toju Duke, Anselm Levskaya, Sanjay Ghemawat, Sunipa Dev, Henryk Michalewski, Xavier Garcia, Vedant Misra, Kevin Robinson, Liam Fedus, Denny Zhou, Daphne Ippolito, David Luan, Hyeontaek Lim, Barret Zoph, Alexander Spiridonov, Ryan Sepassi, David Dohan, Shivani Agrawal, Mark Omernick, Andrew M. Dai, Thanumalayan Sankaranarayanan Pillai, Marie Pellat, Aitor Lewkowycz, Erica Moreira, Rewon Child, Oleksandr Polozov, Katherine Lee, Zongwei Zhou, Xuezhi Wang, Brennan Saeta, Mark Diaz, Orhan Firat, Michele Catasta, Jason Wei, Kathy Meier-Hellstern, Douglas Eck, Jeff Dean, Slav Petrov, and Noah Fiedel. PaLM: Scaling language modeling with pathways, 2022. arXiv:2204.02311.
- Dan Hendrycks and Kevin Gimpel. Bridging nonlinearities and stochastic regularizers with gaussian error linear units, 2016. arXiv:1606.08415.
- Stefan Elfving, Eiji Uchibe, and Kenji Doya. Sigmoid-weighted linear units for neural network function approximation in reinforcement learning, 2017. URL <https://arxiv.org/abs/1702.03118>.
- Yann N. Dauphin, Angela Fan, Michael Auli, and David Grangier. Language modeling with gated convolutional networks, 2017. URL <https://arxiv.org/abs/1612.08083>.
- Noam Shazeer. GLU variants improve transformer, 2020. arXiv:2002.05202.
- Jianlin Su, Yu Lu, Shengfeng Pan, Bo Wen, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding, 2021.

- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *Proc. of ICLR*, 2015.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *Proc. of ICLR*, 2019.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Nee-lakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In *Proc. of NeurIPS*, 2020.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models, 2020. arXiv:2001.08361.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katie Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. Training compute-optimal large language models, 2022. arXiv:2203.15556.
- Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *Proc. of ICLR*, 2020.
- Yao-Hung Hubert Tsai, Shaojie Bai, Makoto Yamada, Louis-Philippe Morency, and Ruslan Salakhutdinov. Transformer dissection: An unified understanding for transformer’s attention via the lens of kernel. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4344–4353, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1443. URL <https://aclanthology.org/D19-1443/>.
- Amirhossein Kazemnejad, Inkit Padhi, Karthikeyan Natesan, Payel Das, and Siva Reddy. The impact of positional encoding on length generalization in transformers. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=Dr12gcjz1>.